

AI Foundations | AIFo

Summary

CONTENTS

1. Machine Learning (ML)	2
2. Probability	2
2.1. Random Variables	2
2.2. Discrete random variables	2
2.2.1. Multiple Random Variables	3
2.3. Conditional Probability	3
2.4. Continuous random variables	4
2.5. Distributions	6
2.5.1. Univariate normal distribution	6
2.5.2. Bernoulli distribution	7
3. Tree Diagrams and sequential events	8
3.1. Vocab for human experiments	10
4. Linear Regression	10
4.1. Ingredients of Machine Learning	10
4.1.1. The 4 main ones	10
4.1.2. Others	10
4.2. Linear Regression	11
4.2.1. Residual sum of square	11
4.2.2. Model	12
4.2.3. Mean squared error	12
4.2.4. Correlation and causality	12
4.2.5. Multiple linear regression	13

1. MACHINE LEARNING (ML)

There are many different AI research fields and AI applications. Algorithms and applications where a computer learns from data are typically in the field of statistical machine learning (ML).

Computer: performs a mapping from input to output.

Processing: any sort of non-trivial “calculation” (mapping) or information processing

Input: any sort of data/information (sensory input, bits, mechanical configuration...)

Output: any sort of response (data, actions, new state of a system...)

Data Visualization: Gives intuitive understanding of structure in the data, helps in identifying patterns, detecting outliers and data quality. Raw data → Information

TODO:

plot examples

<https://ourworldindata.org/>

<https://informationisbeautiful.net/>

2. PROBABILITY

2.1. RANDOM VARIABLES

A random variable X is a function from the sample space to the real numbers.

$$X : S \rightarrow \mathbb{R}$$

Random Variables come in two flavours:

discrete: X takes any of a finite or countable set of values, e.g. $\{-8, 1.5, 1.5, 2.693, 5, 6.3, 10\}$

continuous: X takes any value of an uncountable range, e.g. the real numbers in the interval $(2, 7)$.

$\Pr(X = x)$ is the probability that the random variable X takes the value x . It's often written as $P(x)$ or $p(x)$ or $P_X(x)$ or $\mathbb{P}(x)$.

TODO:

MathFML explained this a lot better

Definition 1: Range

The range of a random variable X , shown by $\text{Range}(X)$ or R_X , is the set of possible values of X .

2.2. DISCRETE RANDOM VARIABLES

The Probability Mass Function (PMF) of a discrete random variable is a function $P(x)$ that provides the probability for each value x of a discrete random variable X .

TODO:

$$0 \leq P_X(x) \leq 1$$

$$\sum_{x \in R_X} P_X(x) = 1$$

$$\text{for any set } A \subset R_X, P(X \in A) = \sum_{x \in A} P_X(x)$$

Example 1:

Dice rolls

Value x of the random Variable X	1	2	3	4	5	6
$P(X = x)$	$1/6$	$1/6$	$1/6$	$1/6$	$1/6$	$1/6$

From the PMF, we can calculate the expected value

$$E[X] = \sum_{i=1}^{\infty} P(x_i) \cdot x_i$$

2.2.1. Multiple Random Variables

The (joint) properties of multiple random variables are defined by the **Joint Probability Mass Function**, usually simply called Joint Probability.

For **independent** random variables, the joint probability is simply the product of the individual probabilities

$$P(X \cap Y) = P(X) \cdot P(Y)$$

Example 2: Rolling two dice

The random experiment of rolling die 1 and die 2 together has 36 possible outcomes (events). The Joint Probability is:

$$P(X = 5, Y = 4) = \frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$$

	$X = 1$	$X = 2$	$X = 3$	$X = 4$	$X = 5$	$X = 6$
$Y = 1$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$
$Y = 2$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$
$Y = 3$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$
$Y = 4$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$
$Y = 5$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$
$Y = 6$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$	$1/36$

2.3. CONDITIONAL PROBABILITY

For **dependent** random variables we use the conditional probability of X given Y defined as

$$P(X|Y) = \frac{P(X \cap Y)}{P(Y)}, \text{ when } P(Y) > 0$$

TODO:

$$P(X \cap Y) = P(X|Y) \cdot P(Y)$$

$$X \cap Y = \emptyset \Rightarrow P(X|Y) = 0$$

$$Y \subset X \Rightarrow P(X|Y) = 1$$

$$X \subset Y \Rightarrow P(X|Y) = \frac{P(X)}{P(Y)}$$

$$P(X \cap Y) = P(X)P(Y|X) = P(Y)P(X|Y)$$

Independence: Two events X and Y are independent iff $P(X \cap Y) = P(X)P(Y)$

Law of Total Probability: If $Y_1, Y_2, Y_3, \dots$ is a partition of the sample space S , then for any event X we have

$$P(X) = \sum_i P(X \cap Y_i) = \sum_i P(X|Y_i)P(Y_i)$$

Bayes' Rule: For any two Events X and Y , where $P(X) \neq 0$, we have

$$P(Y|X) = \frac{P(X|Y) \cdot P(Y)}{P(X)}$$

$$\text{Interpretation: Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Normalizer}}$$

Conditional independence: Two events X and Y are conditionally independent given an event Z with $P(Z) > 0$ if

$$P(X \cap Y|Z) = P(X|Z)P(Y|Z) \\ \Rightarrow P(X|Y \cap Z) = P(X|Z)$$

Example 3: Chance of rain

Let X be the event to observe clouds (0=no clouds, 1=small clouds, 2=big clouds) and Y the event that it rains (0=no rain, 1=light rain, 2=moderate rain, 3=heavy rain)

TODO:

Given a known joint distribution of two discrete random variables, say, X and Y , the marginal distribution of either variable – X for example – is the probability distribution of X when the values of Y are not taken into consideration. This can be calculated by summing the joint probability distribution over all values of Y .

	$X = 0$	$X = 1$	$X = 2$
$Y = 0$	0.35	0.21	0.03
$Y = 1$	0.10	0.07	0.04
$Y = 2$	0.00	0.05	0.05
$Y = 3$	0.00	0.02	0.08

$$P(X) = \sum_Y P(X, Y)$$

$$P(Y) = \sum_X P(X, Y)$$

2.4. CONTINUOUS RANDOM VARIABLES

For a continuous random variable X following holds:

$$P(-\infty < x < \infty) = 1$$

$$P(a < x < b) = \int_a^b f(x) dx$$

$$P(x) = \emptyset$$

The **uniform probability distribution** $\text{unif}(0, 1)$ characterizes an experiment in which one number is chosen from the sample space $\Omega = [0; 1]$ in such a way, that all numbers of Ω have an equal chance to occur.

Whenever $\Omega \subset \mathbb{R}$, we can use $\mathbb{P}$ to define the **cumulative distribution function** (CDF)

$$\begin{aligned} F(\alpha) &= \mathbb{P}((-\infty; \alpha] \cap \Omega) \\ \lim_{\alpha \rightarrow -\infty} F(\alpha) &= 0 \\ \lim_{\alpha \rightarrow \infty} F(\alpha) &= 1 \end{aligned}$$

A function $f : \mathbb{R} \rightarrow \mathbb{R}^+$, such that

$$\int_{-\infty}^{\infty} f(t) dt = 1$$

is called **probability density function** (PDF). With any PDF we can associate a CDF

$$F(\alpha) = \int_{-\infty}^{\alpha} f(t) dt$$

and a probability measure $\mathbb{P}$ that associates the probability of the event $E \subset \mathbb{R}$ to occur with the area of all points underneath the function $f(t)$ whose t -values reside in E :


$$\mathbb{P}(E) = \int_{t \in E} f(t) dt$$

Example 4: Cumulative distribution function F and probability density function f for $\text{unif}(a, b)$

$$F : \begin{cases} \mathbb{R} \rightarrow [0; 1] \\ x \mapsto \mathbb{1}_{[a; b]}(x) \cdot \frac{x-a}{b-a} + \mathbb{1}_{(b; \infty)}(x) \end{cases}$$

$$F'(x) = f(x) = \begin{cases} \frac{1}{b-a} & \text{if } a < x < b \\ 0 & \text{else} \end{cases} = \frac{\mathbb{1}_{(a; b)}(x)}{b-a} = \text{density for } X \sim \text{unif}(a, b)$$

$$\mathbb{P}(X \leq x) = \int_{-\infty}^x f(t) dt = \int_{-\infty}^x \frac{\mathbb{1}_{(a; b)}(t)}{b-a} dt \stackrel{a < x < b}{=} \int_a^x \frac{1}{b-a} dt = \frac{x-a}{b-a}$$



2.5. DISTRIBUTIONS


2.5.1. Univariate normal distribution

The probability density of the normal / Gaussian distribution is given by

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

in which μ and σ are parameters, denoted as **mean value** (μ) and **standard deviation** (σ).

$$\int_{\mu-\sigma}^{\mu+\sigma} \frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{1}{2\sigma^2}(x-\mu)^2} dt \approx 0.68$$



The value of the CDF evaluated at x represents the area of the PDF from $-\infty$ to x , thus giving the probability of an experiment be in $(-\infty; x]$.

All normal distributions can be defined with the help of the so called **standard normal distribution**

$$\varphi(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}$$

for which $\mu = 0$ and $\sigma = 1$. The general normal distribution can then be expressed in terms of the standard normal distribution using the formula

$$f(x|\mu, \sigma) = \frac{1}{\sigma} \varphi\left(\frac{x - \mu}{\sigma}\right)$$

The **error function** is defined through

$$\text{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$$

or its Taylor series

$$\text{erf}(x) = \frac{2}{\sqrt{\pi}} \sum_{k=0}^{\infty} \frac{(-1)^k}{k!(2k+1)} x^{2k+1}$$

and can be approximated through hyperbolic functions

$$\text{erf}(x) \approx \tanh\left(\frac{2}{\sqrt{\pi}}\left(x + \frac{11}{123}x^3\right)\right)$$

Given this function, one can show, that the CDF of the normal distribution is given by


$$F(x|\mu, \sigma) = \frac{1}{2} \left(1 + \operatorname{erf} \left(\frac{x - \mu}{\sqrt{2}\sigma} \right) \right)$$

The **68–95–99.7 rule**, also known as the empirical rule, for a normal distribution is a shorthand used to remember the percentage of values that lie within an interval estimate in a normal distribution: approximately 68%, 95%, and 99.7% of the values lie within one, two, and three standard deviations of the mean, respectively.

$$\mathbb{P}(\mu - 1\sigma \leq X \leq \mu + 1\sigma) \approx 68.27\%$$

$$\mathbb{P}(\mu - 2\sigma \leq X \leq \mu + 2\sigma) \approx 95.45\%$$

$$\mathbb{P}(\mu - 3\sigma \leq X \leq \mu + 3\sigma) \approx 99.73\%$$



2.5.2. Bernoulli distribution

A random variable X is said to be a Bernoulli random variable with parameter p , shown as $X \sim \text{Bernoulli}(p)$, if its PMF is given by

$$P_X(x) = \begin{cases} p & \text{for } x = 1 \\ 1 - p & \text{for } x = 0 \\ 0 & \text{otherwise} \end{cases}$$

where $0 < p < 1$

A Bernoulli random variable is associated with a certain event A . If event A occurs, then $X = 1$; otherwise $X = 0$. For this reason the Bernoulli random variable is also called the indicator random variable. The indicator random variable I_A for an event A is defined by

$$I_A = \begin{cases} 1 & \text{if } A \text{ occurs} \\ 0 & \text{otherwise} \end{cases}$$

The indicator random variable for an event A has Bernoulli distribution with parameter $p = P(A)$, so we can write $I_A \sim \text{Bernoulli}(P(A))$.

TODO:

https://www.probabilitycourse.com/chapter3/3_1_5_special_discrete_distr.php

3. TREE DIAGRAMS AND SEQUENTIAL EVENTS

A **probabilistic model** that explains how data is generated (**generative model**). Generative Models connect domain knowledge with data.


The 2-step model captures the causality of observable events.

- Step 1 lists all our **hypotheses** (mutually exclusive, exhaustive) that could have caused an outcome. (Example: U_1, U_2)
- Step 2 is an enumeration of all possible observable **consequences** (Example: W, S)

Example 5: Urns: Tree Diagram

There are two urns, U_1 and U_2 . Urn U_1 contains 5 silver and 5 white balls, and urn U_2 contains 9 silver and 1 white ball. The probability of drawing a ball from U_1 or U_2 is equal (50/50). What is the probability of drawing a white ball?

We can visualize the probabilities using a **tree diagram**:



The probability of drawing a white ball is therefore 30%, assuming that each urn is chosen with equal probability.

We can use Bayes' theorem to "go backwards" and start with the observation and deduce the cause.

$$P(\text{Hypothesis}|\text{Evidence}) = \frac{P(\text{Evidence}|\text{Hypothesis}) \cdot P(\text{Hypothesis})}{P(\text{Evidence})}$$

Bayesian inference is the statistical method that uses Bayes' theorem to update the probability of a hypothesis after seeing new data. It starts with a **prior probability** (belief before the evidence) and combines it with a likelihood to produce a **posterior probability** (updated belief).

Example 6: Urns: Bayesian inference

We have seen that the choice of urns yield following probabilities (prior):

$$P(U_1) = 0.5$$

$$P(U_2) = 0.5$$

If we observe the drawing of a white ball, we can deduce which urn was more likely to be chosen from (posterior):

$$P(U_1|W) = \frac{P(W|U_1) \cdot P(U_1)}{P(W)} = \frac{0.5 \cdot 0.5}{0.3} = 0.8\bar{3}$$

$$P(U_2|W) = \frac{P(W|U_2) \cdot P(U_2)}{P(W)} = \frac{0.1 \cdot 0.5}{0.3} = 0.1\bar{6}$$

Thus the **prior distribution** of equally distributed probabilities changes to the **posterior distribution** favoring U_1 when provided with the observed data.

Given multiple observed results we can use **Recursive Bayesian estimation** (also called **Bayesian filtering**), which is a general probabilistic approach for estimating an unknown PDF recursively.

$$p(x_0, \dots, x_k, z_1, \dots, z_k) = p(x_0) \prod_{i=1}^k p(z_i|x_i)p(x_i|x_{i-1})$$

Example 7: Urns: Recursive Bayesian estimation

Given the observed outcomes of white, silver, white ($\{W, S, W\}$), we can construct our posterior distribution recursively. We start with the prior $P(U_1) = 0.5$ and $P(U_2) = 0.5$.

1) After the first white ball (W):

$$P(U_1|W) \approx 0.833$$

$$P(U_2|W) \approx 0.167$$

2) Using these as the new priors, we observe a silver ball (S):

$$P(U_1|W, S) = \frac{P(S|U_1) \cdot 0.833}{P(S)} = \frac{0.5 \cdot 0.833}{0.5 \cdot 0.833 + 0.9 \cdot 0.167} \approx 0.73$$

$$P(U_2|W, S) = \frac{P(S|U_2) \cdot 0.167}{P(S)} = \frac{0.9 \cdot 0.167}{0.5 \cdot 0.833 + 0.9 \cdot 0.167} \approx 0.27$$

3) Using these as the new priors, we observe another white ball (W):

$$P(U_1|W, S, W) = \frac{0.5 \cdot 0.73}{0.5 \cdot 0.73 + 0.1 \cdot 0.27} \approx 0.93$$

$$P(U_2|W, S, W) = \frac{0.1 \cdot 0.27}{0.5 \cdot 0.73 + 0.1 \cdot 0.27} \approx 0.07$$

The evidence of two white balls strongly overwhelms the single silver ball, significantly increasing our confidence that the balls have been drawn from U_1 .

TODO:

Kahneman & Tversky

3.1. VOCAB FOR HUMAN EXPERIMENTS

Prevalence: Proportion of a population who have a specific characteristic in a given time period.

Specificity: Percentage of people who test negative for a specific condition among a group of people who do not have the condition. No test is 100% specific because some people who do not have the condition will test positive for it (false positive).

Sensitivity: How well a test can detect a specific condition in people who actually have condition. No test has 100% sensitivity because some people who have the condition will not be identified by the test (false-negative test result).

4. LINEAR REGRESSION

Linear Models are the most simple model to explain a relationship between “Input” and “Output”. Linear Regression is a standard method to find an optimal linear model

4.1. INGREDIENTS OF MACHINE LEARNING

4.1.1. The 4 main ones

Data: The dataset we are given plus the pre-processing pipe-line including cleansing, feature-engineering, data- augmentation etc.

Cost-Function (Loss): A formal (mathematical) expression for “good” and “bad”. Mean Squared Error (MSE) is commonly used.

Model: As simple as a two parameter, linear model $\hat{y}_i = mx_i + b$ or as complicated as a million-parameter Neural Network. Different tasks require different models (e.g. regression, decision-tree, ...)

Optimization Procedure: An algorithm that changes the parameters of the model such that the cost-function is minimized. (e.g. Stochastic Gradient Descent (SGD), ADAM, ...)

4.1.2. Others

Sounds like generic slop ngl

DevOps/MLOps Versioning. Reproducibility: Building efficient pipe-lines is difficult. Following tool-specific recommendations and reference- implementations helps.

Visualization and evaluation of the learning Process: Learning curves, Performance measures, Tensorboard

Cross-Validation & Regularization: The goal is to train models that generalize well to unseen data (and get an estimate of the generalization error)

Baselines and Performance Monitoring: Implement a trivial baseline model. Every advanced model you train should beat it. Define what “better” means, and keep track of the performance evaluations. Deployed models need monitoring.

TODO:

CySec confusion matrix, ground truth, model prediction

TODO:

W4 S7 Different Machine Learning Techniques and applications

4.2. LINEAR REGRESSION

Linear Regression is a simple method to analyse data and quantify a relationship between variables. It has countless applications, most of them fall in one of two categories:

Interpretation: We want to understand if some input has an effect on the output. E.g. Is there a relationship between smoking cigarettes and the risk of lung cancer?

Prediction: Given a new “x” (e.g. Smoking rate), use the model to predict (estimate) the “y” (e.g. Death rate)

In ML, linear regression falls into the category of **supervised learning**: we are given both input data (X) and labels (Y). The ML algorithm learns a linear relationship between X and Y.

4.2.1. Residual sum of square


Residual sum of square (RSS) is a statistical method that helps identify the level of discrepancy in a dataset not predicted by a regression model. Thus, it measures the variance in the value of the observed data when compared to its predicted value as per the regression model.

$$RSS = \sum_{i=1}^N \underbrace{(y_i - f(x_i))^2}_{\substack{\text{Represented as red} \\ \text{squares in the example}}}$$

$$RSS(m, b) = \sum_{i=1}^N (y_i - (mx_i + b))^2 \geq 0, RSS : \mathbb{R}^2 \rightarrow \mathbb{R}$$

Example 8: Linear regression of salaries by age

$$RSS(m, b) = \sum_{i=1}^N ((mx_i + b) - y_i)^2 \geq 0, RSS : \mathbb{R}^2 \rightarrow \mathbb{R}$$



4.2.2. Model

In ML, we use the term **model** for any mathematical function that “explains the data”. More formally:

$$y_i \approx f(x_i)$$

$$y_i = f(x_i) + \varepsilon_i$$

Where ε_i is “unexplained noise”. It is often assumed that ε_i follows a normal distribution.

The function f can be anything from a constant up to a multi-million parameter deep neural network. The goal of ML is to find the model which explains the data as good as possible. Instead of approximating y_i , we calculate an estimate $\hat{y}_i$ of the (usually unknown) y_i

$$\hat{y}_i \approx f(x_i)$$

In linear regression, we decide to only consider a linear relationship between the inputs and outputs. That is, before starting any learning algorithm, we constrain the model space to the family of linear functions. Such a design decision imposes an inductive bias on the model

In the simplest case, m and b are scalars and the linear model therefore has only two free parameters. The goal is to identify m and b for which the linear model “best explains the data”

$$\hat{y}_i = mx_i + b$$

m is usually called **slope**, b the **intercept** (aka **bias**). Often w for “weights” is used: $\hat{y}_i = w_1x_i + w_0$

TODO:

Is y a linear function of x in this model?

$$y = ax^2 + bx + c$$

No. But the model is linear in the parameters a, b, c . Linear regression only needs this second property.

We can apply linear regression to arbitrarily complex models, as long as the model is linear in the unknowns.

4.2.3. Mean squared error

Mean squared error (MSE) is the average of the squares of the errors. It is usually divided by 2.

$$\text{MSE} = \frac{\text{RSS}}{2N} = \frac{1}{2N} \sum_{i=1}^N e_i^2$$

where

$$\hat{y}_i = mx_i + b$$

$$e_i = y_i - \hat{y}_i$$

and e_i (the difference) is called the **residual**.

4.2.4. Correlation and causality

Correlation refers to the degree to which a pair of variables are linearly related. It can be quantified using the **Pearson correlation coefficient**. Showing correlation is not hard, convincingly identifying causality on the other hand is.


4.2.4.1. Pearson correlation coefficient

The Pearson correlation coefficient (PCC) is a correlation coefficient that measures linear correlation between two sets of data. It is the ratio between the covariance of two variables and the product of their

standard deviations. Thus, it is essentially a normalized measurement of the covariance, such that the result always has a value between -1 and 1 .

$$\text{PCC}_{X,Y} = \frac{\sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^N (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^N (y_i - \bar{y})^2}}$$

- 1 shows a perfect positive correlation where both variables increase / decrease together at a constant rate
- -1 shows a perfect negative correlation where one variable increases as the other decreases proportionally
- 0 shows no linear relationship, meaning changes in one variable do not predict changes in the other



4.2.5. Multiple linear regression

Same concept but with more factors and weights:

$$\hat{y}_i = \beta_1 x_{1i} + \beta_2 x_{2i} + \dots + \beta_N x_{Ni} + \alpha_i$$

where

- y_i is an observed/measured quantity. Example: blood pressure
- $x_{1i} \dots x_{Ni}$ are “factors”. Examples: x_{1i} : age, x_{2i} : weight, x_{3i} : sex, x_{4i} : drug_dosage, ...
- $\beta_1 \dots \beta_N$ are weights. How much does each factor x explain the outcome y ?

4.2.5.1. Matrix notation

The dataset can be written as N points (x, y) , where x is a vector with p features.

$$\hat{y} = X\beta + \alpha$$

where

$$X = \begin{pmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \dots & x_{Np} \end{pmatrix}, \quad \beta = \begin{pmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_N \end{pmatrix}, \quad y = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix}$$

TODO:

W4 S18,20