

youkined from <https://github.com/Nicicalu/OST-CN2-Spik>

OSPF (IGP)
Equal Cost Multi-Path (ECMP) load balancing, fast convergence, widely used.
Isos IP Port 89 on L3. Uses 224.0.0.5 (all routers), 224.0.0.6 (all DR,BDR)
DR Election: Highest IF prio → if tie, highest Router ID. Priority 0 = ineligible for DR/BDR. No preemption. BDR = same rules, second best candidate
Virtual links: Cannot go through more than one area, can only run through standard non-backbone areas. (eg. not over stubby areas)

ROUTER OPERATION
– Establish neighbor adjacencies and exchange LSAs
– Build the Link-State Database (LSDB)
– Run Dijkstra's SPF algorithm:
– Intra-area change: SPF recalculation
– Inter-area change: No SPF needed → ABR handles updates
– Build the routing table from SPF results

ROUTE TYPES
O: Intra-area: Local area → Local area. Type 1 and 2 LSAs
B: Build the Link-State Database (LSDB) for the OSPF domain.
O E1 or O E2: External: Local area → External (eg. BGP), from ASBR
O E1: Cost = internal + external metrics
O E2: Cost = external metrics. Default external route for OSPF
Cost calculation: Reference Bandwidth(def=100 Mbps)/Interface Bandwidth

AREAS
AS Split into areas with 32-bit Area ID (e.g., 0.0.0.0 = Area 0)
Backbone Area (Area 0): The OSPF domain of the backbone must connect to all other areas (directly/virtual links), must be contiguous (no disjointed segments), should not contain end-user networks, summarizes topology of other areas
Non-Backbone Areas: Connect end users and local resources, All inter-area traffic must transit the backbone

ROUTER TYPES
Area Border Router (ABR): Connects two or more OSPF areas, must have 1 interface in backbone. 1 OSPF DB per Area
Internal Router (IR): All interfaces belong to the same area (non-backbone)
Backbone Router (BR): At least one interface in Area 0, includes ABRs and routers internal to the backbone

AS Boundary Router (ASBR): Connects to external AS/Network (e.g., BGP). Advertises external routes into OSPF, can be in backbone or non-backbone

DESIGN RULES
Rule 1: Backbone (Area 0) must be contiguous → no partitions allowed
Rule 2: Every non-backbone area must connect to Area 0

LINK STATE ADVERTISEMENT (LSA) TYPES
(1) **Router-LSA:** All routers list directly connected links (so all outgoing interfaces with state and cost) (intra-area)
(2) **Network-LSA:** DR lists all routers and DR in broadcast/multi-access network (intra-area)
(3) **Summary-LSA:** ABR advertises networks from other areas, flooded in all areas (that are not "totally stubby" (inter-area)
(4) **ASBR-Summary:** ABR advertises path to ASBRs. ASBR sends Type 1 to external bit, ABR converts to Type 4 and floods (inter-area)
(5) **AS-External-LSA:** ASBR/ABR advertises external routes (e.g., BGP), flooded to all non-stub areas → only normal areas (inter-area)
(7) **NSAA External:** Like Type 5 but inside NSAA. Converted to Type 7 by ABR. Shown in routing table as [O N1 or O N2]

AREA TYPES (ALLOWED LSA TYPES)
Stubby = Blocks external LSA [No 5], Totally = Blocks inter-area LSA [No 3/4]
Default: (e.g., 0.0.0.0) from [DR ABR], Cannot contain ASBR [No ASBR]
Standard (1,2,3,5): Normal
Stub Area (1,2,3,5): [No 4/5] [DR ABR] [No ASBR]
Totally Stubby Area (1,2,3,5): [No 3/4/5] [DR ABR] [No ASBR]
Not-So-Stubby Area (1,2,3,7): [No 4/5], no DR, allows 1 ASBR [LSA 7 → 5]
Totally NSAA (1,2,7): [No 3/4/5] [DR ABR] [LSA 7 → 5]

PACKET TYPES
(1) **Hello:** Usually Multicast. Discover, maintain, and verify neighbors. Forms adjacencies. Election of DR/BDR in broadcast networks. Contains network mask (sending routers' IF), Hello interval (p2p,broadcast: def=10s), Options, Priority (election), Router dead time (def=40s), DR/BDR IP, Neighbors.
(2) **Database Description (DBD/DD):** Unicast. Exchange summaries of LSAs during adjacency formation (headers only). Contains IF Max. MTU, Options, I/M/M/S bits (Initial, More, Master-slave bit), DD Seq. Num., LSA Header
(3) **Link State Request (LSR):** Unicast. Request specific LSAs. Contains Link State Type (router/net), Link State ID, Advertising Router (sender address)
(4) **Link State Update (LSU):** Multi/Unicast. Implicit ack. Flood new/updated LSAs. Answer to LSR (unicast). Contains Nr. of LSAs, full LSAs information.
(5) **Link State Acknowledgment (LSAck):** Multi/Unicast. Explicit ack. Confirms receipt of LSAs for reliable flooding. Acks may be grouped together.

SUB-PROTOCOLS
Hello Protocol: Neighbor discovery. Parameter negotiation. Maintain logical adjacencies on P2P, P2MP, virtual links. Elet DR/BDR on broadcast and NBMA networks. Continuously sends hello packets to maintain bidirectional connectivity, failure to receive = neighbor down (RouterDeadInterval)
Database Sync Protocol: Sync LSDB using DBD packets.

NEIGHBOR STATES
Down: Initial state, no Hello packets exchanged or after RouterDeadInterval
Init: Received Hello packet from neighbor but without its own Router ID yet
2-Way: Seen its own Router ID. DR/BDR election is taking place in NBMA
ExStart: (DBD) Bi-dir comm; highest Router-ID=master. Determine init Seq Nr
Exchange: Exchange of DBD packets (LSA headers)
Loading: Missing LSAs are requested. (LSRLSU)
Full: Databases are synchronized. Implicit ack. Flood new/updated LSAs.
OSPF ROUTING AND ECMP
– Each router runs Dijkstra per area; link cost = metric from LSAs (1–65535)
– Perform more specific match (CIDR), else O > O > E1 > N1 > E2 > N2
– Routes added to RIB/FIB based on computed next hops
– ECMP: Modified Dijkstra supports Equal-Cost Multi-Path if multiple paths have same cost → routes added with multiple next-hops for load balancing

IS-IS (IGP)
Widely used by ISPs, Fast convergence, ECMP, Extendable ES: End-host devices are called End Systems (ES)
IS: Routers are called Intermediate Systems (IS)
CONNECTIONLESS NETWORK SERVICES (CLNS)
LNS: ISO Layer 3 datagram service; supports CLNP, ES-IS, IS-IS.
CLNP: Connectionless Network Protocol, similar to IP, used in ISO stack
IS-IS: host → router discovery protocol
IS-IS: router ↔ router. Link-state routing protocol
Integrated IS-IS: Allows IP routing with IS-IS

NETWORK SERVICE ACCESS POINT (NSAP)
Variable in length (up to 20 bytes), Identifies 49.001.1. 0000.0000.0003. 00
node, not IF. Example from IP: 192.168.1.24 49.0001.1921.6800.1024. 00
Area ID (L1) System ID (R3) NSEL

Network Entity Title (NET): Unique router identifier. NSEL=0, Starts with 49
Authority & Format identifier (AFI): NSAP format-authority that assigned it
Initial Domain Identifier (IDI): Variable len, identifies administrative domain
Domain Specific Part Format identifier (DFI): Specifies format
NSEL (N-Selector): always 0n for routers (1b)
System ID: Unique per router (often based on Io-IP/MAC) (6b)
Area ID: Used for routing hierarchy (like OSPF areas) (AFI+DI+DSP)

PACKET TYPES (ALL IN L1 OR L2)
All PDUs contains Header with shared common fields + specific routing-related information (Type, Length and Value (TLV)) used to extend protocol
IIH (Hello): Build/maintain adjacencies. Includes system ID, holding time, prio
Functions: discover, build, maintain
Interval: 10s (default); DIS sends every 3.3s on LANS
Multiplier: Missed Hello time! → Holdtime = Interval × Multiplier(def=3)
Multicast: [L1 01-80-c2-00-00-14] [L2 01-80-c2-00-00-15]
LSP (Link State PDU): Contains topology info including prefixes with costs; Unicast on p2p or flooded throughout the area (OSPF: LSA Type 1)
Complete SNRP: Sent by DIS, lists all known LSPs (OSPF: DBD)
PSNP (Partial SNRP): Used to request missing or acknowledge received LSPs

ROUTER TYPES
L1 Router: Within one area; all have same LSPDB. No inter-area routing
L2 Router: Backbone router; all have same LSPDB, Routes between areas
L1/L2 Router: Acts as both; separates LSPDBs, redistributes between levels

ADJACENCIES
Backbone must be contiguous: L1 ↔ L1-L2 ↔ L2
L1: Are addresses match unless configured otherwise
L2: Alongside L1, unless router is L1 only, areas must not match

POINT-TO-POINT (NB1)
Initialized by receipt of IS-Is, exchange p2p IIHs padded to MTU
CSNP: Sent once at adjacency startup
LSP: Advertises topology changes (link-state info)
PSNP: Acknowledges received LSPs or requests missing ones
BROADCAST/MULTI-ACCESS (OIS REQUIRED)
Not triggered by IS-Is, instead broadcast IIHs with all neighbors MAC's
Designated Intermediate System (DIS): Create, update pseudonode LSPs.
Floods LSPs. Similar to DR in OSPF. No backup DIS.
DIS Election: Highest priority (0–127) (Cisco def=84) = Highest SNPA (MAC)
Preemption: Enabled → higher prio router automatically takes over DIS Role
Pseudonode: Carried out by DIS. Multiaccess links. Separate for L1,L2
Passive interfaces: Advertise network prefixes without adjacency forming

PATH SELECTION
Default interface metric = 10. **Lower Metric better:** L1 intra-area > L2 intra-area > Leaked L2 → L1 (internal metric) > L1 external (external metric) > L2 external (external metric) > Leaked L2 → L1 (external metric)

LEVEL 1 ROUTING
– Intra-area routing only (L1 Area like OSPF Totally Stubby Area)
– L1 routers install a default route to nearest L1-L2 for inter-area traffic
– L1-L2: Do not advertise L2 routes into L1 area (unless route leaking active)
– L1-L2: Set Attached bit to signal L2 connectivity to backbone
– **Distribution Bit:** Set to 1 on L2 → L1; blocks re-advertisement L1 → L2.
– **Route-Leaking:** injects a more specific route into L1 to improve routing

LEVEL 2 ROUTING
– Routing between areas (inter-area)
– L1,L2 routers inject L1,L2 routes into L2 topology
– L1 routes are redistributed into L2 with L1 metric preserved in L2 LSP

IS-IS VS OSPF

Feature	IS-IS	OSPF
Layer	L2 (CLNS)	L3 (IP, proto 89)
Encapsulation	No IP, uses TLVs	IP packets
Hello Type	IIH	Hello packet
Area Model	L1/L2	Backbone + Areas
Metric	Cost (default 10)	Cost (bandwidth)
Router ID	System ID (6B)	32-bit Router ID
Adj. Types	L1, L2, L1/L2	DR/BDR, P2P
LSDB	Per level (L1/L2)	Per area
Scaling	Large-scale ISP core	Enterprise/campus
Routing info	TLVs (flexible)	Fixed LSA types
Type	Link-State	Link-State
Shortest Path Calculation	Dijkstra	Dijkstra

BGP (EGP)
Path-vector routing protocol. Stable, flexible, scalable. Based on TCP
IGBP
Routers in same AS, AD=200, more trusted (lower security overhead). AS-Path and Next-hop not modified

SCALABILITY
– iBGP does not re-advertise routes between iBGP peers → full mesh required (caused no loop prevention)
– Session count = $n \times (n-1) / 2$ (e.g., 5 routers = 10 sessions, 10 = 45)
ROUTE REFLECTORS (RR)
Solves iBGP full-mesh scaling by allowing selective route reflection. Clients only peer with RR, unaware they're clients. Only the RR needs special config. From non-client: RR advertises to clients only
From client: RR advertises to all clients + non-clients
From eBGP peer: RR advertises to all clients + non-clients

EBGP
Routers in different ASs, AD=20, stricter policy enforcement. Prepends ASN to AS-Path, modifies next-hop IP. Fails if own ASN in AS-Path. TTL=1
AS-Override: Allows route of same ASN across different customer sites (e.g., Swisscom); rewrites ASN to avoid loop detection

CONFIG
next-hop-self fixing iBGP: neighbor (neighbor IP) next-hop-self (overriding)
AUTONOMOUS SYSTEM NUMBERS (ASN)
– Unique ID for each AS; required for Internet routing with BGP
– Private: [Legacy 16-bit 64512–65535] [32-bit 4200000000–4294967294]
PEERING / NEIGHBORS
– Two routers with a BGP TCP session (port 179) are called peers/neighbors
– Each BGP router is a **BGP speaker**
– Supports CIDR, route aggregation; decisions based on policies/rules

PATH ATTRIBUTES
Used for route control and policy enforcement
Well-known mandatory: Always present (e.g., AS-Path, Origin, Next Hop)
Well-known discretionary: Optional but recognized by all (e.g., Local Pref)
Optional transitive: Passed between ASes (e.g., Community, Aggregator)
Optional non-transitive: Not passed across ASes (e.g., MED, Weight)
NLRI: Routing table info: prefix, prefix length, and associated path attributes

MESSAGES
OPEN: Establishes session; includes version, ASN, Hold Time, BGP Identifier, optional params. **Hold Time:** Heartbeat in seconds (Cisco def=180s), reset by KEEPALIVE/UPDATE; 0 = session down. **BGP Identifier:** 32-bit router ID, manually set or highest loopback/active IP; used for loop prevention
KEEPALIVE: Sent every 1/3 of Hold Time; ensures neighbor liveness
UPDATE: Advertises new routes, withdraws old ones, or both; includes NLRI (prefix + path attributes); can act as KEEPALIVE
NOTIFICATION: Sent on session error (e.g. hold timer expired); terminates session immediately

NETWORK STATEMENTS
Advertise a specific network prefix to BGP peers (only the best is advertised, even if multiple exist). Must exist in the RIB so that it's imported in BGP table
Configure: next-hop: 0.0.0.0, origin attribute: i, weight: 32768 (Cisco)
Static/RP: next-hop: NH in RIB, MED: IGP metric, rest same as above

BEST PATH CALCULATION
– BGP maintains all received paths per prefix but advertises only the best one
– Best path is installed in RIB, recalculated on:
– Next-hop reachability change
– Interface failure to eBGP peer
– Redistribution change
– New/withdrawn path received
– Influence:
– Outbound BGP policy → inbound traffic behavior
– Inbound BGP policy → outbound traffic behavior

BEST PATH SELECTION (CISCO ORDER):
(1) Prefer highest **Weight** (Cisco-specific, local to router)
(2) Prefer highest **Local Preference** (global within AS)
(3) Prefer routes **originated by the router** (only small i path, NH 0.0.0.0)
(4) Prefer **shortest AS path** (only length is compared)
(5) Prefer **lowest origin type** (IGP < EGP < Incomplete (I on origin))
(6) Prefer **lowest MED** (Multi-Ext Discriminator) (also called metric)
(7) Prefer **external (EBGP) over internal (IBGP)**
(8) For iBGP: prefer path with **lowest IGP metric** to next-hop
(9) For eBGP: Prefer **lowest** (more stable) path
(10) **Next-highest BGP router ID**
(11) Prefer path from **lowest neighbor IP** address

ROUTE FILTERING
– Filters control which routes are received/advertised
– Used for security, traffic shaping, memory optimization
– Tools: **prefix-list** (IP), **filter-list** (AS-path), **route-map** (flexible match/set)

COMMUNITIES
32b optional, transitive tag (e.g. ASN: valUe, 65000:100). Used to mark routes for policy control across ASs. Can be added, modified, or removed at each hop

PEERING VS. TRANSIT
Transit: ISP provides full reachability (paid relationship)
Peering: ISP exchanges selected routes; equal relationship, usually unpaid
AS Path filtering: Used to avoid getting transit: ip as-path access-list 10 permit * \$

INTERNET EXCHANGE POINT (IXP)
– Facility where networks exchange traffic via BGP peering
– Reduces transit costs, latency, and offloads upstream links

PUBLIC PEERING
– Members peer via shared switch fabric and a **route server**
– Route server distributes routes but stays out of data path (NEXT_HOP unchanged)
– Minimal policy control; one BGP session to route server
– Simplified setup (one legal contract)

PRIVATE PEERING
– Direct BGP sessions between two parties (1:1)
– May use public or private interconnects
– Full policy control per neighbor
– Requires one session and legal contract per peer

Examples: Equinix, SwissIX (non-profit)
ENTERPRISE CONNECTIVITY OPTIONS
SingleDual: denotes how many links there are
Multi-Homed-/Homed: denotes how many ISPs are connected
Single-Homed: One ISP, one link (BGP or static). Simple but no redundancy
Dual-Homed: One ISP, two links (routers). Redundancy within same provider
Multi-Homed: Multiple ISPs, improved redundancy and routing control, but avoid being a transit – advertise only customer-owned prefixes

Dual-Multihomed: (Multi-Homed) but two links per ISP. Most redundancy
TRAFFIC ENGINEERING (TE) - INBOUND VS OUTBOUND
Local Pref: [OUT] [highest T] Lower local pref on backup path
Med: [IN] [lowest T] Higher MED on backup path
AS-Path Prepending: [IN] [shortest T] Add multiple own ASN on backup path
TE Limitation: AS controls [OUT], [IN] control limited - ISP may ignore

TE with Aggregate: Prefer primary IPN with summarized routes; advertise specific prefixes on backup for fallover

INTERNET
Transit: One ISP provides reachability to all destinations
Peering: ISPs provide each other reachability to portions of their routing table

RESOURCE PUBLIC KEY INFRASTRUCTURE (RPKI)
– Framework to validate that prefix is legitimately originated by a specific ASN
– Prevents route hijacking and accidental mis-originations (route leaks)
– RPKI validates origin only – not the AS-path
KEY REQUIREMENTS
Trust Anchor (TA): Root CAs of the 5 RIRs; issue certs for resource holders
ROA – Route Origin Authorization: Digitally signed object, authorizes ASNs to announce prefix (AS, prefix that AS can originate, max prefix length)
RPKI Validators: downloads+verifies+stores Validated ROA Payloads (VRPs)

RPKI VALIDATION STATES
Valid: Prefix + ASN match ROA
Invalid: Prefix found, but ASN mismatch or prefix too long
Not Found / Unknown: No matching ROA
VALIDATION DETAILS
– Use TALLs (Trust Anchor Locators) to fetch data from RIRs
– Validate cryptographic signatures (via X.509 certs with SHA 3779)
– Outputs VRPs; invalid objects are discarded
– Update frequency: >=24h (recommended), ~30–60min (practice)
– RRDp (RFC 8182) replacing rsync (uses HTTPS)

MONITORING
Event tracking, BGP hijack detection, Route leak detection, RPKI status check, Reachability tracking, AS path change tracking, AS path visualization

NETWORK DESIGN
Pillars: Scalability, Speed, Availability, Security, Manageability → overall Cost
AVAILABILITY CONCEPTS (MOST IMPORTANT REQUIREMENT)
MTBF: Mean Time Between Failures, **MTTR:** Mean Time to Repair
$$M = \frac{MTBF}{MTBF + MTTR}$$

$$M_{combined} = \left(\sum_{n=1}^N \frac{1}{MTBF_n} \right)^{-1}$$
 parallel
– Lower MTTR + higher MTBF = more availability
REDUNDANCY
– Adds reliability, decreases MTBF but increases MTTR and complexity
– Balance: resilience vs. manageability

Backup Paths: Duplicate devices/links on primary path, build extra links for redundancy, consider backup link capacity, consider failover speed
Load Balancing: ECMP, EtherChannel, Port-Channel
HIERARCHICAL DESIGN
Access: Connect end devices, high port count, port security, L2, QoS marking
Distribution: L3, policy control, HSRP/VRP, loop protection, small fault domain
Core: High-speed backbone, L3 only, no policy, scalable/redundant, no security

Composed Core: Combines Core + Distribution (small/medium networks)
FLAT DESIGN
– Any-to-any; small networks, MPLS/LAN setups
– Lacks scalability and control

FABRIC DESIGN
Modern design using Underlay/Overlay
Underlay: transport (e.g., IP, MPLS)
Overlay: logical virtual topology (e.g., EVPN)
ENTERPRISE CAMPUS
– 100s/1000s of users, multiple buildings, one physical location
– Multiple interconnected LANS, connected via Ethernet and Wireless
FHRP – FIRST HOP REDUNDANCY PROTOCOLS
– Ensures default gateway is always reachable, PCs can only have one
– Enables fast failover during router failure

VRP – Virtual Router Redundancy Protocol (Multivendor): Shared virtual IP, real MACs per router, Master router handles forwarding
HSRP – Hot Standby Router Protocol (Cisco): Shared virtual IP + virtual MAC, One active, others in standby → faster
GLBP – Gateway Load Balancing Protocol (Cisco): Load balancing + redundancy, Shared virtual IP + multiple virtual MACs, Roles: AVF: Answers ARP and sends virtual MAC addresses of AVFs AVF: Forwards traffic

DATA CENTER DESIGNS
ToR - Top of Rack: switches per rack, less cabling, easy expansions/ exchanges per rack; scalable glass fiber, ideal for high service density (full racks). But more switches, more ports, more L2 Srv-2 Srv traffic, more STP to be managed
EoR - End of Row: 1 switch per row, less switches, higher utilization of ports, switches all at one place, better L2 availability between racks. But more cabling

TRAFFIC DIRECTIONS
North-South: Between external networks (client-server, in/out DC)
East-West: Within DC (e.g., server-server, storage)

THREE-TIER DC
– Access – Aggregation – Core (like Hierarchical)
– Optimized for North-South traffic
– Not ideal for East-West communication
LEAF-SPINE ARCHITECTURE
– Two-tier: Leaf switches (access) connect to Spines (core)
– High performance, low latency
– Scalable, ideal for East-West traffic

MULTICAST
Use-cases 1-to-Many: Streaming, software updates, music-on-hold
Use-cases Many-to-Many: Gaming, VR, stock data, group chat
Benefits: Efficient bandwidth, lower server/CPU load, no redundancy
Properties: UDP-based, Apps must handle dups, duplicates, out-of-order
Source: Sends to group IP; doesn't need to join
Receiver: Must explicitly join group to receive traffic
MULTICAST ADDRESS RANGES
224.0.0.0 – 224.0.0.255: Link-local, TTL=1 (not forwarded by routers) (IGMP)
224.0.1.0 – 224.0.1.255: Reserved by IANA, routable (NTP)
232.0.0.0 – 232.255.255.255: Source-Specific Multicast (SSM)
233.0.0.0 – 238.255.255.255: Globally scoped multicast (routable)
239.0.0.0 – 239.255.255.255: Administratively scoped (private space)

BROADCAST BASICS
L3 routes between subnets; L2 floods within same subnet
L2 (Flooding): MAC ffff.ffff.ffff floods flood to all ports in VLAN
L3 (Bridging): 255.255.255.255: local broadcast, not routed. Directed broadcast for specific subnet (e.g. 18.1.1.255): can be routed if enabled
L2 MULTICAST: MAC MAPPING
(1) Get IP: Example multicast IP address: 239.5.5.5
(2) Convert to Binary: 239.5.5.5 = 11101111.00000101.00000101.00000101
→ Take only the last 23 bits: 00000101.00000101.00000101
(3) Map to MAC Prefix: Fixed 0100.5E → Map last 23 bits to: 0100.5E05.0505

INTERNET GROUP MANAGEMENT PROTOCOL (IGMP)
Purpose: Manages group membership for IPv4 multicast on each segment
IGMPv1: Basic join via query-response mechanism, no explicit group leave. Router sends **membership queries** every 60s to 224.0.0.1, removes group after timeout if no report received. Receiver doesn't know who the source is
IGMPv2: Adds **Leave Group** message. Supports **Group-Specific Queries** (e.g. when someone leaves group (anyone still interested in group xy?), reducing broadcast overhead. Receiver still doesn't know who the source is)
IGMPv3: Adds **source filtering** (Include/Exclude lists). Enables **Source-Specific Multicast (SSM)** – receiver receives traffic only from selected source(s), no need for Rendezvous Point (RP) anymore. Adds support for application-level access control and filtering. Can also be used in ASM (Any Source Multicast), but mainly with SSM. (224.0.0.22 all IGMPv3 routers)

IGMP SNOOPING
Without snooping: multicast = broadcast on VLAN
With snooping: switch listens to IGMP messages & builds a forwarding table
Default: snooping is enabled; switch needs IGMP Query to operate
REVERSE PATH FORWARDING (RPF)
To avoid loops, verifies that a multicast packet arrives on the interface that a unicast packet destined for the multicast source would be forwarded out of.
RENDEZVOUS POINT (RP)
– Used in Shared Tree (*,G) setups with PIM-SM
– RP acts as the common meeting point for sources and receivers
– RPF check is performed toward the RP (not the source)
– Once the source is known, routers may switch to a Source Tree (S,G)
SHARED TREE VS. SOURCE-BASED TREE
Shared Tree (*,G):
– IGMP host sends a membership report (IGMP Join)
– Router adds (*,G) entry to multicast routing table
– * means 'any source' → source is unknown/unspecified
Source-Based Tree (S,G):
– Built when router receives an (S,G) Join/report from IGMP host
– S = known multicast source, G = group
– Router adds (S,G) to mroute table once source is known
PROTOCOL INDEPENDENT MULTICAST (PIM)
Only between routers. Relies entirely on the unicast routing table (RIB) for multicast forwarding decisions. Uses Reverse Path Forwarding (RPF), Protocol-independent: static routes, OSPF, IS-IS, etc.
PIM-DM – DENSE MODE (NO RP)
Push Model: Floods multicast to all interfaces; prunes where no receivers
(1) **Flooding:** Source sends traffic to multicast-enabled links using unicast RIB
(2) **Distribution Tree:** Initially entire network (shared tree rooted at source)
(3) **Prune Mechanism:** Routers without interested receivers send prunes upstream to remove themselves from the tree
(4) **State Maintenance:** Routers track source, receivers, interfaces to forward/prune per group
PIM-SM – SPARSE MODE (ASM, SSM)
Pull Model: Multicast traffic is only sent where requested. Works with IGMP to detect interested receivers and uses unicast routing for forwarding.
(1) **Join/Prune:** Routers send explicit join/prune messages to request or stop receiving multicast for group (G) to other routers
(2) **Forwarding:** Routers only forward multicast packets for group (G) on interfaces from which explicit joins were received
ASM – ANY SOURCE MULTICAST (HOW IT WORKS WITH RP)
– Works with IGMPv1 or IGMPv2 → receiver does not know the source
– Receiver sends **IGMP Join (*,G)** to its first-hop router
– First-hop router forwards **PIM Join (*,G)** hop-by-hop toward the Rendezvous Point (RP)
– RP acts as a common meeting point for sources and receivers
– Sources send multicast traffic to the RP via a **PIM Register** tunnel
– All routers in the multicast domain must know the RP location
SSM – SOURCE SPECIFIC MULTICAST
– Receiver subscribes using IGMPv3, providing both source (s) and group (G) to the first-hop router
– No Rendezvous Point (RP) required → **PIM-SSM builds only (s,G)** Short-est Path Trees (SPT)
– **No shared tree (*,G)** mode used; SSM is source-directed
– IANA reserved 232.0.0.0/8 for SSM in IPv4
– Join messages are forwarded hop-by-hop toward the source to establish forwarding path
– Uses unicast routing table (RPF) to maintain loop-free delivery
PTN SPARSE-DENSE MODE
PIM Sparse Mode: Pull model – multicast traffic is forwarded only on request
– **Join:** User sends IGMPv3 request to join a specific group, then pruned
Sparse-Dense Mode: Supports both modes per multicast group; choice depends on RP availability
PIM-DM VS. PIM-SM
PIM is protocol-independent: it relies on the unicast routing table for RPF checks and to forward joins toward the source or Rendezvous Point (RP).
– **PIM Dense Mode (DM)** – "Push" model
– Floods multicast traffic throughout the network
– Prunes back where traffic unwanted
– **PIM Sparse Mode (SM)** – "Pull" model
– Traffic sent only on request
– Requires explicit **Join** messages
Usage Recommendation:
Dense Mode: Best for small or tightly scoped networks where most devices need multicast
Sparse Mode: Preferred for large-scale or distributed environments where multicast receivers are few or spread out
VXLAN
Issues of L2: STP. Max amount of VLANs (4094). Large MAC Address tables
VXLAN (Virtual Extensible LAN): Tunnels Ethernet (Layer 2) over IP using MAC-in-UDP encapsulation (Port 4789). For flexible and scalable network segmentation.
VNID (VXLAN Network Identifier): 24-bit identifier (up to 16 million segments) that defines the VXLAN broadcast domain.
VTEP (Virtual Tunnel Endpoint): Device (switch, router, or host) responsible for encapsulating/decapsulating VXLAN traffic.
NVE (Network Virtual Interface): Logical interface on a VTEP used for VXLAN tunnel operations.
TUNNEL
– VXLAN establishes IP tunnels between VTEPs to extend Layer 2 networks across Layer 3 boundaries.
– VXLAN enables both L2 and L3 VPN functionality in overlay networks.
– VXLAN traffic is encapsulated in UDP (default port: 4789).
FRAME FORMAT
– Ethernet frame → VXLAN Header → UDP → Outer IP Header.
– The VXLAN header contains the 24-bit VNID and flags.
– Outer headers allow Layer 2 frames to traverse IP underlay networks.
VIRTUAL NETWORK IDENTIFIER (VNI)
– 24-bit VXLAN Network Identifier uniquely defines VXLAN segments.
– Replaces traditional VLAN IDs (12-bit), enabling ~16 million logic segments.
– Used by VTEPs to map traffic into corresponding Layer 2 domains.
VXLAN TUNNEL ENDPOINT (VTEP)
– Connects the overlay (VXLAN) and underlay (IP) networks.
– **Types:**
– **Software VTEP:** Located on hypervisors using virtual switches.
– **Hardware VTEP:** Located on routers/switches with ASICs for performance.
– **Interfaces:**

- **VTEP IP Interface:** Connects to the underlay network and handles encapsulation.
- **VNI Interface:** Virtual interface per segment (like SVI); handles segregation of Layer 2 domains.

MAC ADDRESS LEARNING

On control plane: happens proactively
on data plane: ad-hoc with flooding

- Each VTEP maintains a VXLAN mapping table linking destination MAC addresses to remote VTEP IPs.
- **Learning via ARP:**
 - Host H1 sends ARP request, switches learn H1's MAC.
 - ARP request is flooded to H2.
 - H2 responds; switches learn H2's MAC.
- **Learning Methods:**
 - **Static VXLAN:** Manual MAC-to-VTEP mappings. Doesn't scale well; BUM traffic is inefficient.
 - **Multicast VXLAN:** VTEPs join multicast groups per VNI. Scales better, offloads BUM replication. 20+ VTEPs = there is too much traffic, doesn't scale well
 - **MP-BGP EVPN:** Modern solution using BGP as control plane. Dynamically learns MAC/IP info.

- **EVPN**
- Overcome flood-and-learn limitations, doesn't rely on data plane learning, utilizes robust control plane MP-BGP, works with different encapsulation tech-niques (VXLAN, MPLS), excellent scalability, I2 and I3 Support.

MP-BGP EVPN (MULTIPROTOCOL BGP FOR ETHERNET VPN)

- Enables protocol-based VTEP discovery and host reachability via control-plane learning
- Reduces flooding by replacing data-plane learning
- Extends BGP with multiprotocol capabilities (AFI/SAFI)
- Uses **MP_REACH_NLRI** and **MP_UNREACH_NLRI** for route advertise-ment and withdrawal

EVPN ROUTE TYPES

- Type 2 – Host Advertisement:** Advertises host MAC (mandatory), optionally IP, along with L2VNI and optionally L3VNI. Used for MAC learning, ARP suppression, and host mobility. Sent when host connects to VTEP.
- Type 5 – Subnet Advertisement:** Advertises IP prefix + prefix length with L3VNI. Used for inter-subnet routing. VTEP redistributes connected/static/dynamic IP routes. Additional attributes: L3VNI, extended communities.

HOST DELETION & MOVE

Host Deletion: When a host detaches, its ARP (default: 1500s) and MAC entry (default: 1800s) time out on the VTEP. Upon aging, the VTEP withdraws the host's MAC/L2VNI and IP/L3VNI advertisements.

Host Move: When a host moves to a new VTEP, the new VTEP advertises updated reachability with a higher move sequence number. The old VTEP withdraws its entry, completing the migration.

ROUTE DISTINGUISHER (RD) VS. ROUTE TARGET (RT)

- **Route Distinguisher (RD):** Uniquely identifies VPN routes – allows same IP prefix to be used in different VPNs. Can be IPv4 or ASN used to make routes unique in BGP (VPNv4/v6). Forms VPNv4 NLRI: **RD:IPv4 prefix**
- **Route Target (RT):** Controls route import/export between VRFs. *Used as extended BGP community.*

How RTs Work:

- A route is tagged with an RT when advertised by BGP.
- Other VRFs import the route if the RT matches their import policy.
- Allows overlapping or shared connectivity between tenants (e.g., shared services).
- **Format:** Typically in the form *ASN:nn* or *IP-nn*, e.g., 65009:100, 1:118
- **Multiple RTs can be used:** A route can have multiple RTs for flexible policies (e.g., one RT for VPN, another for shared services)

EVPN (ETHERNET VPN) – L2

KEY FEATURES

- L2 bridging across L3 networks

BGP Control Plane: Distributes MAC info (no flooding)
VXLAN Overlay: Encapsulates L2 in L3 UDP (data plane)
Multi-Tenancy: via VNI segmentation
Redundancy: All-active multihoming, ECMP, fast convergence

USE CASES

- Multi-tenant datacenter interconnects (DCI)
- Extending L2 over WAN between remote sites
- Scalable, segmented L2 fabrics

BGP CONTROL PLANE

- PEs learn MACs from local CEs (data plane)
- MACs advertised via BGP (control plane)
- Uses Route Distinguishers and MPLS labels
- Remote PEs update L2 RIB/FIB with MAC and next-hop info
- Enables seamless L2 across IP/MPLS backbone

EVPN NLRI

- EVPN uses MP-BGP with specific AFI/SAFI
- Supports multiple route types and attributes
- Unsupported routes are dropped by BGP

AUTODISCOVERY VIA ROUTE REFLECTORS

- Route Reflector (RR) avoids full-mesh iBGP
- RR reflects EVPN routes to other PEs
- RR doesn't participate in EVPN or pseudowires
- RR needs only *address-family l2vpn evpn*
- L2VPN RIB stores endpoint/VFI info for control plane
- **BGP_UPDATE** from spines contain **ORIGINATOR_ID** (origin leaf)

HOST DETECTION

- Host connects to VTEP → MAC learned locally
- VTEP advertises MAC + L2VNI via BGP EVPN
- MAC learning follows normal Ethernet semantics

INGRESS REPLICATION (IR)

- BUM traffic, when Multicast underlay network is not used, handle multi-destination traffic (ARP → unicast)

EARLY ARP TERMINATION (ARP SUPPRESSION)

- Avoids flooding ARP requests
- VTEP queries control plane for MAC/IP/VNI mapping
- If known → direct unicast (no broadcast)

SILENT HOST FLOW (FALLBACK)

- If IP/MAC unknown → ARP sent via **ingress replication**
- Replicated ARP request goes to remote VTEPs
- Only correct host responds → update reflected to all VTEPs
- Future traffic uses updated BGP mapping

VRF – VIRTUAL ROUTING AND FORWARDING

- Multiple isolated routing tables on one device
- Each tenant = one VRF → traffic isolation
- Supports independent policies per tenant
- Key for scaling and multi-customer separation

IRB – INTEGRATED ROUTING AND BRIDGING

- Enables inter-VLAN routing inside EVPN
- Avoids central gateway → 'no traffic tromboning'
- Two modes: Symmetric and Asymmetric

SYMMETRIC IRB (L2 + L3)

- Routing/bridging on ingress + egress VTEPs
- Uses L3 Transit VNI (same in both directions), one L3 VNI per VRF (Tenant)
- Scales well; clean separation of MAC and IP

ASYMMETRIC IRB (L2)

- Routing only on ingress, bridging on egress
- VXLAN uses destination VNI in both directions
- one L2 VNI per VLAN/Subnet
- Simple config, but requires all VLANs/VNIs on all VTEPs

DISTRIBUTED ANYCAST GATEWAY (DAG)

- Same gateway IP+MAC on all VTEPs
- Enables local default gateway for hosts
- Supports mobility + optimal forwarding

L3 HOST DETECTION

- Host sends ARP/ND to local VTEP
- VTEP learns MAC/L2VNI and IP/L3VNI
- Info is advertised in EVPN (control plane)

MPLS

- Label Switched Path (LSP) → pre-determined path across MPLS network
- advantage eBGP between PE-CE: No mutual redistribution, same routing process
- encrypt traffic flowing over MPLS L3VPN backbone? yes (e.g. bank)
- Unicast Reverse Path Forwarding (uRPF): checks source of each packet & verifies that source is in routing table
- control plane (e.g. OSPF) → to learn labels
- iBGP used to exchange NLRI (RD, RT, IPv4 Prefix, NextHop &VPN Label) between PE
- imp-null = networks are directly connected, no more label switching

WAN

- Connects remote LANs via SPs for data/voice/video; key needs: bandwidth, control, design, resilience, mgmt. **Requirements:**
- **Bandwidth:** App needs, peak usage, reserve for VoIP
- **Control/Security:** Trust provider? No full control
- **Availability:** Redundancy, SLA for failures
- **Mgmt:** Inband vs out-of-band

PRIVATE WAN

- **Point-to-Point:** Leased L2 line (Ethernet); monthly fee; private circuit
- **Dark Fiber:** Physical fiber lease; costly; ISPs prefer selling lambdas
- **Connection-oriented:** Predefined path, packets carry IDs (ATM, Frame Relay)
- **Connectionless:** No setup; full address in each packet (Ethernet, MPLS VPN)

TERMINOLOGY

- **CE - Customer Edge:** no knowledge of MPLS, no labels; connected to PE
- **PE - Provider Edge:** connected to CE; runs iBGP and LDP; uses VRFs
- **P - Provider or LSR (Label Switch Router):** inside MPLS VPN, no CE connection; forwards labels

DATABASES, PLANES

RIB (Routing IB (Information Base)): Learned prefixes from routing protocols

FIB (Forwarding IB): Built from RIB; only best routes for forwarding
LIB (Label IB): All label mappings: 1 label per prefix
LFIB (Label Forwarding IB): Built from LIB; used for actual forwarding decisions. (LFIB only contains currently best LSP (decision): Routing Protocol)

Control Plane: Builds routing/label tables (RIB, LIB)

Data Plane: Forwards packets (FIB, LFIB); pushes/swaps/pops labels (see below)

MPLS HEADER

4-byte header before IP:

- **Label (20b)** – actual MPLS label
- **EXP (3b)** – QoS/CoS, Now called Traffic Class (TC)
- **S-bit (1b)** – bottom of label stack indicator, 1 = True = last label before IP header
- **TTL (8b)** – time-to-live (eq tual IP TTL)

TTL MPLS

- ingress PE router decrements IP TTL field & copies packet's IP TTL field into new MPLS TTL
- P routers decrements MPLS TTL
- egress PE router decrements MPLS TTL, pops final MPLS header, copies IP TTL
- traceroute receive ICMP Time Exceeded, Provider doesn't want to expose MPLS network to fire: disable MPLS TTL propagation (on PE), PE set MPLS TTL = 255, egress leaves PE original IP TTL unchanged
- = MPLS network appears as single router hop from TTL perspective

LABEL DISTRIBUTION PROTOCOL (LDP) – CONTROL PLANE

- Distributes labels to neighbors using control plane
- **Hello messages:** Sent via UDP (Port 646) to 224.0.0.2 to discover neighbors
- TCP (Port 646) connection is used to exchange label bindings (prefix to local label)
- Routers advertise all local bindings after TCP session is up
- Label mapping used to build LIB → LFIB
- LDP router ID must be reachable (via routing table)
- Each router manages local labels independently

MPLS L3 DATA PLANE

- VPN traffic uses 2 labels (stacked):

- Outer label:** Transport label (LDP); identifies LSP between ingress/egress PE
- Inner label:** VPN label (MP-BGP); identifies customer VRF
- Push:** Ingress PE; classify and label packets
- Swap:** P router; replaces label, forwards based on new label
- Pop:** Egress PE removes label; sends original packet to CE
- Penultimate Hop Popping:** MPLS feature, penultimate router removes the outer MPLS label before forwarding to egress PE, default enabled, ISPs disable it

VRF TABLES

- VRF = Virtual Routing and Forwarding table (Virtual router inside a PE.
- Maintains isolated RIB + FIB per customer.)
- Stores separate routing info per customer (VRF attached customer)
- Exists per MPLS-aware PE router, one per attached customer

- Contains: RIB, FIB, and separate routing process per CE

VPNv4

- 64-bit RD + 32-bit IPv4 = 96-bit VPNv4 prefix
- transferring VPNv4 between PE router → Multiprotocol iBGP (MP-iBGP)

OVERLAY TECHNOLOGIES

MODERN PROVIDER NETWORK

- MPLS:** Label-based forwarding (fast, scalable)
- LDP:** Distributes labels for MPLS paths
- IGP:** Underlay routing (e.g., OSPF, IS-IS)
- MP-BGP:** Extends BGP to carry VPNv4/v6, EVPN routes

DRAWBACKS OF TRADITIONAL NETWORKS

- Control Plane:** LDP/RSVP-TE adds complexity
- Scalability:** Per-flow/path state limits growth; LSP and signaling overhead increase rapidly

- OAM:**
 - **Troubleshooting:** Traceroute less useful in MPLS; labels hide topology
 - **Traffic Eng.:** LDP lacks TE; relies only on IGP cost
- Fast Reroute:** Limited coverage; microloops possible

SEGMENT ROUTING (SR)

Source routing: Sender defines full path using Segment List (Segment = instruction)

SID = Segment Identifier: Each SID is 1 instruction (e.g., forward via ECMP, specific iface, or to a service)

State in packet: No per-flow state in network; intermediate nodes follow SID instructions

No new control plane: Uses existing protocols (OSPF, IS-IS, BGP) with extensions; no LDP or RSVP-TE needed

Segment List: Ordered SID list carried in packet header; defines full route

Simple but powerful: Enables TE, fast reroute, policy routing

SEGMENT LIST OPERATIONS

Push: Insert SIDs into packet; set active SID (top of list)

Continue: Active SID not yet completed; keep processing it

Next: Current SID completed; activate next SID in list

SEGMENT SIGNIFICANCE

Global Segments: Known and supported by all SR nodes in the domain, Installed in forwarding tables across the network (e.g. "Forward packet according to shortest path to Node1")

Local Segments: Defined and installed only on originating node, Not forwarded by others, but must be understood network-wide (e.g. "Forward packet on interface to Node2")

Global segments: Defined in the SR Global Block (SRGB) and should be consistent across all nodes; **local segments** are defined in the SR Local Block (SRLB) and are specific to the local SR node.

SR CONTROL PLANE SEGMENT TYPES

IGP Prefix Segment: Global SID tied to IGP prefix (multi-hop); all nodes install forwarding entries

IGP Node Segment: Global SID for a specific node (shortest-path forwarding)

IGP Anycast Segment: Global SID for a group of nodes; traffic sent to nearest

IGP Adjacency Segment: Local SID; direct link to neighbor

L2 Adjacency SID: Local SID for Layer-2 segment (e.g., Ethernet link)

Combining Segments:

- End-to-end paths can mix IGP and BGP segments
 - Traffic to BGP Anycast → more ECMP in data centers
- **SR-MPLS**
 - Reuses existing MPLS data plane – no hardware change needed
 - Segments = MPLS labels; Segment List = label stack (top = active)
 - Segments distributed via IGP/BGP; no LDP required (interoperable if needed)
 - Supports both IPv4 and IPv6 networks

BENEFITS OF SEGMENT ROUTING

Benefits: Simplification (removes protocols, simple operations, admin and mgmt), enhanced Traffic eng. (Delay, Bandwidth, Packet Loss, TE metric, Controller, Source-Node), Seamless deployment, Robust, Network Innovation (zB Container Networking)

Source Router: Balances distributed intelligence with centralized optimization

TI-LFA: Fast reroute technique; protects against link/node failure with micro-loop avoidance and no pre-calculation dependency

Traffic Engineering (TE): Optimizes network performance by analyzing and controlling data flow to reduce congestion and improve QoS

Service Function Chaining (SFC): Chains SDN services in order; automates traffic between VNFs and optimizes routing for performance

QoS

Internet is best effort: no guarantees, no QoS; all traffic treated equally (net neutrality); simple, scalable, but no delivery/order assurance or prioritization

QOE & ROUTE PINNING

QoE – Quality of Experience: Perceived service quality from user perspective

Route Pinning: Keeps flow on a fixed path to prevent oscillation (don't switch immediately to "better" path)

NETWORK PERFORMANCE METRICS

Latency / Delay [ms]: Time for packets to travel src → dest (Voip < 150ms)

End-to-End Delay: Total time sender to receiver

One-Way Delay: From first bit sent to last bit received

Delay Components: Transmission delay (time to push onto link), Processing delay (lookup, queuing), Propagation delay (physical travel time)

Jitter [ms]: Variation in delay between packets, caused by re-routing/queuing (Voip<30ms), Calc: no queue - queued delay

Throughput: Rate of successfully delivered data

Packet Loss [P]: Dropped packets due to congestion or errors (Voip < 1%)

Bandwidth [Gbit/s]: Maximum transfer capacity of a link

QUEUEING ALGORITHMS

FIFO (First-In First-Out): Basic, no prioritization

Priority Queueing (PQ): Multiple queues, serve highest first; others may starve

Round-Robin: One packet per queue in turn (fair, but ignores priority)

Weighted Fair Queueing (WFQ): Round-Robin with weights, e.g., 2 packets from Q1, 4 from Q2

Class-Based WFQ (CBWFQ): WFQ with user-defined classes, queue limits, max bandwidth guaranteed or max % of bandwidth (logical queues based on IP Precedence only)

Low Latency Queueing (LLQ): Adds strict priority queue (priority class) to CBWFQ for delay-sensitive traffic (e.g. voice) (based on IP Precedence, DSCP, src, port, protocol...)

QUEUE MANAGEMENT

Tail Drop: Drops packets when queue full; huge interruption of traffic → same as no connectivity

TCP Global Sync: Many TCP flows back off and restart simultaneously → link underutilization

TCP Starvation: TCP slows down after queue, UDP doesn't → queues filled with UDP, TCP squeezed out

RED: Random early drops before full queue to prevent global sync and TCP collapse. Dropped TCP segments cause TCP sessions to reduce their windows sizes

WRED: RED + DSCP/EXP-based drop logic, prioritizes higher-marked traffic

DSCP / EXP: DSCP (6-bit in IP header) marks packets for QoS; used in Diff-Serv for classifying traffic. EXP (3-bit in MPLS label) serves same purpose within MPLS networks; often mapped from DSCP.

POLICING VS. SHAPING

Policing (Inbound mostly): Drops packets that exceed configured rate limits

Shaping (Outbound): Buffers packets to smooth traffic bursts and conform to profile

QOS MODELS

Best Effort: No guarantees, all traffic treated equally (follows Internet neutral-ity)

Integrated Services (IntServ): End-to-end QoS, per-flow resource reservation, precise but not scalable (uses RSVP)

Differentiated Services (DiffServ): Class-based, scalable approach using marking (e.g.,), no hard guarantees

TRAFFIC MARKING

L3 Marking: ToS byte → DSCP (6 bits) + IP Precedence (3 bits)

L2 Marking: Dot1q header → 802.1p CoS bits

MODULAR QoS CLT (MQC)

Class Map: Define traffic classes (e.g., match voice or video)

Policy Map: Define actions for each class (e.g., limit, shape, priority)

Service Policy: Apply policies to interfaces or directions (in/out)

CDN

Origin Server: Content content source (original files), usually in a datacenter

Edge / CDN Server (POP - Point of Presence): Geographically distributed, caches content

DNS Infrastructure: Directs users to optimal edge server (e.g. via Geo-Routing)

KEY BENEFITS

- **Latency Reduction:** Nearby edge servers reduce round-trip time
- **Availability:** Failover and redundancy in case of node failure
- **Scalability:** Handles traffic spikes via load balancing
- **Cost Optimization:** Reduces backend and transit load on origin
- **DDoS Protection:** Edge servers absorb attacks → not all traffic on one server
- **Global Load Reduction:** Less long-distance traffic across the Internet

REQUEST ROUTING TECHNIQUES

- Decides which edge server should serve a client request
- Goal: Best performance (e.g. proximity, load, responsiveness)

DNS-BASED GEO-ROUTING

- Each edge has a unique IP
 - DNS server picks closest/optimal edge server based on:
 - Resolver IP location (not user!)
 - GeoIP DBs (MaxMind, IP2Location), load, latency, business rules
 - Limitation: DNS Resolver != user location → can cause wrong choice
- **EDNS(0) and CLIENT SUBNET EXTENSION (ECS)**
 - Resolver includes part of client IP in DNS request (e.g. /24 subnet)
 - Authoritative DNS makes better decision based on actual client region
 - Improves accuracy without revealing full IP

ANYCAST WITH BGP

- Same IP (e.g. 7.7.7.7) advertised from multiple locations
- BGP routing decides which path is "best" (AS-path, local pref, etc.)
- No DNS logic or per-client decision → pure BGP convergence

Pros: Fast failover, simple, no app logic needed

Cons: Less control, BGP != best latency, route flapping risk

HTTP CACHING & HEADERS

- Caching is controlled via HTTP headers between clients, proxies, and servers

Cache-Control: Main directive (*no-cache, no-store, max-age, must-reval-idate*, etc.)

Expires: Absolute expiration time (older method, replaced by *Cache-Control*)

ETag: Validator tag (version/hash), used with *If-None-Match*

Last-Modified: Timestamp used with *If-Modified-Since* for revalidation

Age: Time (in seconds) since response was fetched from origin

Validation: Client uses *ETag* or *Last-Modified*; server returns 304 if un-changed

OTHER INFOS

AD: Inter-protocol choice (e.g., OSPF vs RIP) →

Cost/Metric: Intra-protocol choice (e.g., OSPF path A vs B) → lower wins

Routing Preference Order (across protocols):

- Most specific prefix
- Lowest Administrative Distance
- Static default route

Administrative Distances: (Smallest Administra-tive Distance wins)

EVPN BGP ROUTING TABLE INFOS

= > Would not be there if it was L2 VNI BGP Routing Table

Route Distinguisher: 172.16.255.101:32777

Route Type: 2

MAC Address Length: 48

MAC Address: 5254.00f8.29a8

***IP Address Length:** 32

***IP Address:** 10.10.0.100

L2 VNI: 30010

***L3 VNI:** 50000

Remote VTEP IP Address: 172.16.254.101

L2 Route Target: 172.16.255.101:32777

***L3 Route Target:** 65000:50000

TODO'S
OSPF
– passive interfaces
– flooding
– O N1 / O N2 routes
– route summarization
– synchronizing the Isdb
IS-IS
– ES-IS, IS-IS: details
– adjacencies
– route leaking
– summarization
BGP
– comparison to IGP
– Multihop sessions
– NLRI
– Aggregate impact (TE)
– rework
– Best path calculation
INTERNET
– Routing policy
– IXP
AVAILABILITY
– everything
MULTICAST
– Link-local multicast
– IGMP and the querier/Role in Dense Mode/Challenges
TODO:
notes 37+