Mathematical Foundations for Machine Learning | MathFML

Summary

CONTENTS

1. Linear algebra . . . 2
1.1. Standard basis . . . 2
1.2. Vectors . . . 2
1.3. Linear transformations . . . 2
1.3.1. Affine transformations . . . 3
1.4. Matrices . . . 4
1.4.1. Properties . . . 5
1.4.2. Unit matrices . . . 5
1.4.3. Kronecker Delta . . . 6
1.4.4. Matrix product . . . 6
1.5. Tensors . . . 7
2. Functions of several variables . . . 7
2.1. Image classification . . . 7
2.2. Residual sum of square . . . 7
2.3. Visualizing functions . . . 8
2.3.1. Curves and surfaces . . . 10
2.3.2. Vector valued functions of several variables . . . 13
2.4. Calculus of curves . . . 15
2.5. Calculus of real valued functions in many variables . . . 17

1. LINEAR ALGEBRA

Term 1 **: Linear combination**

$$\sum_{i=1}^{k} \lambda_i \vec{v}_i$$

1.1. STANDARD BASIS

Basis vectors $\{\vec{e}_1, \vec{e}_2, \ldots, \vec{e}_n\}$ in $\mathbb{R}^n$ where $\vec{e}_i = \left(0, 0, \ldots, \underbrace{1}_{i\text{th position}}, \ldots, 0\right)^T$

Example 1

$$\left[\begin{pmatrix}1\\0\\0\end{pmatrix}, \begin{pmatrix}0\\1\\0\end{pmatrix}, \begin{pmatrix}0\\0\\1\end{pmatrix}\right] = \mathbb{R}^3$$

It can be shown that any basis of $\mathbb{R}^n$ consists of exactly n linear independent vectors and that conversely any set of n linear independent vectors is a basis of $\mathbb{R}^n$.

1.2. VECTORS

TODO:
group and clean this up

$$\lambda\vec{0} = \vec{0}$$
$$\vec{v} + \vec{0} = \vec{v}$$
$$-\vec{v} = -1 \cdot \vec{v}$$
$$-\vec{v} + \vec{v} = \vec{0}$$
$$(\lambda\mu)\vec{v} = \lambda\left(\mu\vec{v}\right) = \lambda\mu\vec{v}$$
$$\lambda\left(\vec{v} + \vec{w}\right) = \lambda\vec{v} + \lambda\vec{w}$$
$$\vec{v} + \left(\vec{u} + \vec{w}\right) = \left(\vec{v} + \vec{u}\right) + \vec{w} = \vec{v} + \vec{u} + \vec{w}$$

1.3. LINEAR TRANSFORMATIONS

$$T : V \to W$$

Term	***Definition***
Image / range	$\operatorname{im} T = \{T(v) \mid v \in V\} = T(V)$ 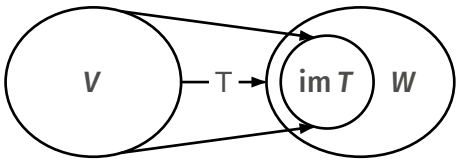

Term	*Definition*
Kernel / nullspace	$\ker T = \{v \in V \mid T(v) = 0\}$ 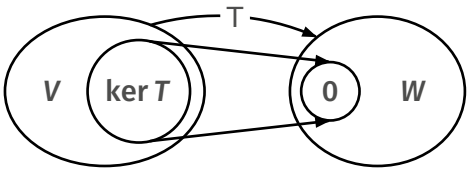

***Observations** 1*

1.1. $\operatorname{im} T \subset W$

1.2. $\ker T \subset V$

1.3. $\vec{0} \in \ker A$

***Example** 2*

$$A \in \mathbb{R}^{n \times n}, T_A : \begin{cases} \mathbb{R}^n \to \mathbb{R}^n \\ \vec{x} \quad \mapsto A\vec{x} \end{cases}$$

$$\ker A = \left\{\vec{x} \in \mathbb{R}^n \mid A\vec{x} = \vec{0}\right\} = \left\{\vec{x} \in \mathbb{R}^n \mid T_A(\vec{x}) = \vec{0}\right\}$$

TODO:

1.3.1. Affine transformations

All linear transformations share the property, that they map the null vector to null:

$$L_A(\vec{0}) = \vec{0}$$

For many applications this is undesirable. Therefore linear transformations are often slightly extended to transformations, which are called **affine transformations** in mathematics. In order to define an affine transformation you need an $r \times c$ matrix M and an r-dimensional vector b, which, in machine learning and statistics, is usually called bias-vector or intercept. The affine transformation is then given by

$$A_{b,M} : \begin{cases} \mathbb{R}^c \to \mathbb{R}^r \\ x \quad \mapsto b + M \cdot x \end{cases}$$

***Example** 3* : **Straight line in one dimension**

$$A_{b,m} : \begin{cases} \mathbb{R} \to \mathbb{R} \\ x \mapsto mx + b \end{cases}$$

1.4. MATRICES

Term	*Definition*
Rank	Maximum number of linearly independent rows or columns $$A \in \mathbb{R}^{m \times n} \Rightarrow \operatorname{rank}(A) = \rho(A) \leq n$$ ***TODO:*** check
Nullity	Number of vectors in the null space $$A \in \mathbb{R}^{m \times n} \Rightarrow \operatorname{nullity}(A) = n - \rho(A) = \dim(\ker A)$$ ***TODO:*** check
Dimension	$$\dim(\ker A) = \operatorname{nullity}(A)$$ $$\vec{x} \in \mathbb{R}^3 \Rightarrow \dim \vec{x} = 3$$
Span	Set of all finite linear combinations of the elements of S of a vector space V. $$\operatorname{span}(S) = \left\{\lambda_1 \vec{v}_1 + \lambda_2 \vec{v}_2 + \ldots + \lambda_n \vec{v}_n \;\middle\vert\; n \in \mathbb{N}, v_1, \ldots, v_n \in S, \lambda_1, \ldots, \lambda_n \in K\right\}$$ $$\operatorname{span}\{(1,0,0),(0,1,0),(0,0,1)\} = \mathbb{R}^3$$
Column space	Span of the column vectors of a matrix $$\operatorname{colsp}(A) = C(A)$$ $$A \in \mathbb{R}^{m \times n}, \dim(\operatorname{colsp}(A)) = n$$
Row space	Span of the row vectors of a matrix $$\operatorname{rowsp}(A) = C\left(A^T\right)$$ $$A \in \mathbb{R}^{m \times n}, \dim(\operatorname{rowsp}(A)) = m$$

***Example** 4*

TODO:

***Observations** 2*

2.1. $\operatorname{rank}(T) + \operatorname{nullity}(T) = \dim V$

2.2. $\dim(\operatorname{im} T) + \dim(\ker T) = \dim(\operatorname{domain} T)$

2.3. $\operatorname{rank}(A) = \dim(\operatorname{rowsp}(A)) = \dim(\operatorname{colsp}(A))$

2.4. $A \in \mathbb{R}^{n \times n} \Rightarrow \dim(\ker A) + \operatorname{rank} A = n$

2.5. $A \in \mathbb{R}^{n \times n}, \ker A = \left\{\vec{0}\right\} \Leftrightarrow \operatorname{rank} A = n$

1.4.1. Properties

Term	*Rules*
Associativity	$(A+B)+C = A+(B+C) = A+B+C$ $(A \cdot B) \cdot C = A \cdot (B \cdot C) = A \cdot B \cdot C$
Distributivity	$C \cdot (A+B) = C \cdot A + C \cdot B$ $(A+B) \cdot C = A \cdot C + B \cdot C$
Commutativity	$A+B = B+A$ $A \cdot B \neq B \cdot A$
Transposing	$(A^T)^T = A$ $(A+B)^T = A^T + B^T$ $(\lambda A)^T = \lambda A^T$ $(A \cdot B)^T = B^T \cdot A^T$ $A_{ij} = A^T_{ji}$
Identity matrix	$\mathbb{1} \cdot A = A \cdot \mathbb{1} = A$ for $A \in \mathbb{R}^{n \times n}$
Inverting	$A \cdot A^{-1} = A^{-1} \cdot A = \mathbb{1}$
Determinate	$\det(\lambda M) = \lambda^n \det(M), M \in \mathbb{R}^{n \times n}$ $\det\begin{pmatrix} A & * \\ 0 & B \end{pmatrix} = \det(A) \cdot \det(B)$ $\det(A \cdot B) = \det(A) \cdot \det(B)$ $\det(A^{-1}) = \frac{1}{\det(A)}$ $\det(A^T) = \det(A)$
Scalars	$(\lambda + \mu)A = \lambda A + \mu A$ $(\lambda\mu)A = \lambda(\mu A)$ $\lambda(B+C) = \lambda B + \lambda C$ $\lambda(BC) = (\lambda B)C = B(\lambda)C$

1.4.2. Unit matrices

Let M be the set of all 2×2 matrices. How could the basis of M look? $\begin{pmatrix} a & b \\ c & d \end{pmatrix}$

Unit matrices:

$$E_{11} = \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix}, E_{12} = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, E_{21} = \begin{pmatrix} 0 & 0 \\ 1 & 0 \end{pmatrix}, E_{22} = \begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix}$$

$$M = \{E_{11}, E_{12}, E_{21}, E_{22}\}$$

$$(E_{11})_{22} = 0, (E_{11})_{11} = 1$$

$$\vec{e}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \in \mathbb{R}^3, (\vec{e}_1)_1 = 1$$

1.4.3. Kronecker Delta

$$\delta_{ij} = \begin{cases} 1\,, i = j \\ 0\,, \text{otherwise} \end{cases}$$

Example 5

$$\delta_{23} = 0, \delta_{11} = 1, \left(\vec{e}_1\right)_j = \delta_{1j} = \begin{cases} 1\,, 1 = j \\ 0\,, \text{otherwise} \end{cases}$$

$$\delta_{m1}\delta_{m2} = 0$$
$$\delta_{kl} - \delta_{lk} = 0$$
$$\left(e_k\right)_i = \delta_{ki} = \delta_{ik}$$
$$\left(E_{kl}\right)_{ij} = \delta_{ki}\delta_{lj}$$
$$\left(E_{rst}\right)_{ijk} = \delta_{ri}\delta_{sj}\delta_{tk}$$

1.4.4. Matrix product

$$A \in \mathbb{R}^{n\times m}, B \in \mathbb{R}^{m\times r}, A \cdot B \in \mathbb{R}^{n\times r}$$

$$A \cdot B = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{i1} & a_{i2} & \dots & a_{im} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nm} \end{pmatrix} \cdot \begin{pmatrix} b_{11} & b_{12} & \dots & b_{1r} \\ \vdots & \vdots & \ddots & \vdots \\ b_{i1} & b_{i2} & \dots & b_{ir} \\ \vdots & \vdots & \ddots & \vdots \\ b_{m1} & b_{m2} & \dots & b_{mr} \end{pmatrix}$$

$$(A \cdot B)_{11} = a_{11}b_{11} + a_{12}b_{12} + \dots + a_{1m}b_{m1} = \sum_{k=1}^{m} a_{1k}b_{k1}$$

$$(A \cdot B)_{ij} = \sum_{k=1}^{m} a_{ik}b_{kj}$$

Shorthand:

$$(A \cdot B)_{ij} = \sum_{k} a_{ik}b_{kj}$$

Transposing:

$$\left[(A \cdot B)^T\right]_{ij} = (A \cdot B)_{ji} = \sum_{k} A_{jk}B_{ik}$$

Example 6 : **Let X be a matrix for which X^{-1} exists. Prove that the inverse of X^T exists and is $\left(X^{-1}\right)^T$**

$$\mathbb{1}^T = \mathbb{1}$$
$$\Rightarrow \left(X^{-1}X\right)^T = \mathbb{1}$$
$$\Leftrightarrow X^T\left(X^{-1}\right)^T = \mathbb{1}$$

1.5. TENSORS

Vectors = tensors of rank 1, matrices = tensors of rank 2

The standart basis for $\mathbb{R}^{n\times m\times l}$ consists of $n \cdot m \cdot l$ basis vectors E_{ijk} where $1 \le i \le n, 1 \le j \le m, 1 \le k \le l$, such that all components of E_{ijk} are zero except at position i, j, k where the value of the component is 1.

2. FUNCTIONS OF SEVERAL VARIABLES

Let D and R be two sets. A mapping $f : D \to R$ that associates to each element $x \in D$ exactly one element of $f(x) \in R$ is called function.

Term	*Definition*
Domain of f	D
Range of f	R
Real valued function	$R \subset \mathbb{R}$
Vector valued function	$R \subset \mathbb{R}^m$
A function of n variables	$D \subset \mathbb{R}^n$

2.1. IMAGE CLASSIFICATION

A 128×256 pixel image I with 3 color channels (RGB) can be represented as a $128 \times 256 \times 3$ matrix (tensor of rank 3). The red value of a pixel at row 100, column 12 can be indexed with $I_{100,12,1}$. A function to determine whether the image was taken at day (0) or night (1) would have the following signature:

$$f : \mathbb{R}^{128\times256\times3} \to \{0, 1\}$$

2.2. RESIDUAL SUM OF SQUARE

Residual sum of square (RSS) is a statistical method that helps identify the level of discrepancy in a dataset not predicted by a regression model. Thus, it measures the variance in the value of the observed data when compared to its predicted value as per the regression model.

$$RSS = \sum_{i=1}^{N} \underbrace{\left(y_i - f(x_i)\right)^2}_{\text{Represented as red squares in the example}}$$

$$RSS(a, b) = \sum_{i=1}^{N} \left(y_i - (mx_i + b)\right)^2 \ge 0, RSS : \mathbb{R}^2 \to \mathbb{R}$$

Example 7 : **Linear regression of salaries by age**

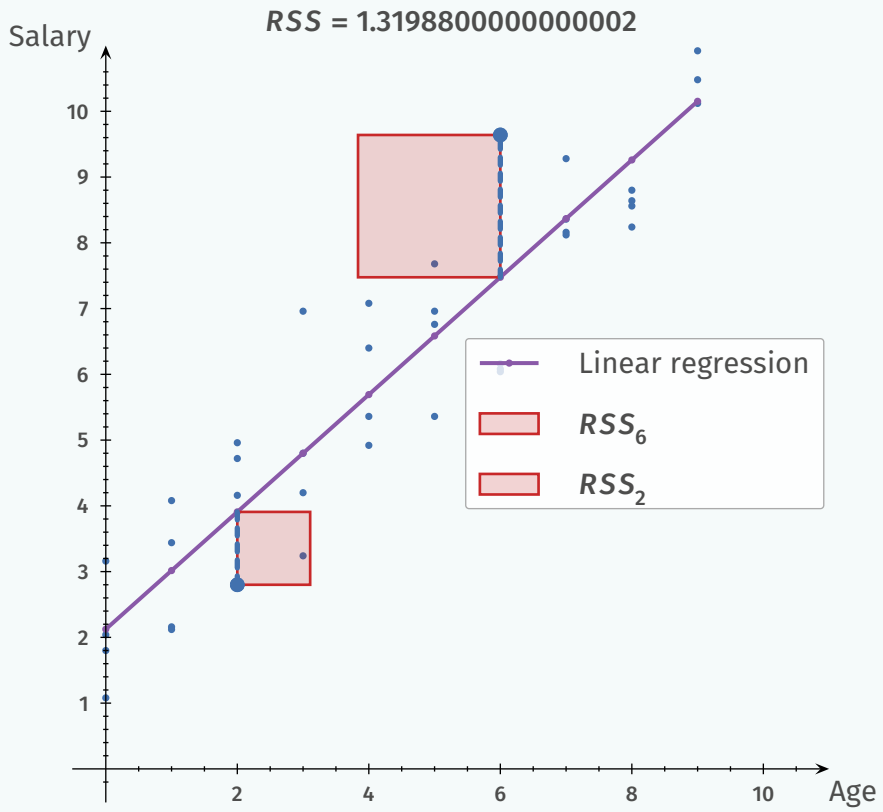

2.3. VISUALIZING FUNCTIONS

If $f : D \to R$ is a function from $D \subset \mathbb{R}^n$ to $R \subset \mathbb{R}^m$ its graph is defined as:

$$\text{graph}(f) = \left\{ (\vec{x}, \vec{y}) \in D \times R \;\middle|\; x \in D \wedge \vec{y} = f(\vec{x}) \right\} \subset D \times R \subset \mathbb{R}^n \times \mathbb{R}^m = \mathbb{R}^{n+m}$$

Due to the fact that the visual imagination of humans is limited to at most 3 dimensions, the graph of a function is a useful concept for graphical illustration if and only if $n + m \leq 3$.

***Example** 8*

The graph of the function

$$f : \begin{cases} \mathbb{R} \to \mathbb{R}^2 \\ t \mapsto (\cos t, \sin t) \end{cases}$$

is

$$\begin{aligned} \text{graph}(f) &= \left\{ (t, x, y) \in \mathbb{R}^3 \;\middle|\; (x, y) = f(t) \wedge t \in \mathbb{R} \right\} \\ &= \left\{ (t, \cos t, \sin t) \in \mathbb{R}^3 \;\middle|\; t \in \mathbb{R} \right\} \end{aligned}$$

It is displayed below and illustrates how f maps t to the corresponding values

TODO:
same diagram as below

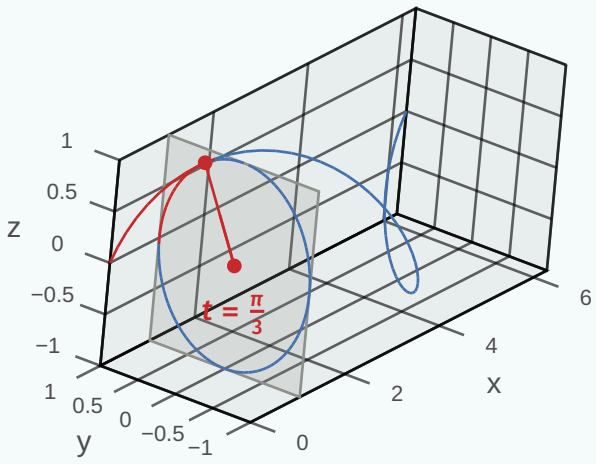

Example 9

The graph of the function

$$g : \begin{cases} \{(x,y) \in \mathbb{R}^2 \mid x^2 + y^2 \leq 1\} \to \mathbb{R}^2 \\ (x,y) \qquad\qquad\qquad\qquad \mapsto \sqrt{1 - x^2 - y^2} \end{cases}$$

is

$$\begin{aligned} \operatorname{graph}(g) &= \left\{ (x,y,z) \in \mathbb{R}^3 \;\middle|\; z = \sqrt{1 - x^2 - y^2} \wedge x^2 + y^2 \leq 1 \right\} \\ &= \left\{ \left(x, y, \sqrt{1 - x^2 - y^2}\right) \in \mathbb{R}^3 \;\middle|\; x^2 + y^2 \leq 1 \right\} \end{aligned}$$

It is displayed below and illustrates how f maps (x,y) to the corresponding values

$$g((x,y)) = \left(x, y, \sqrt{1 - x^2 - y^2}\right)$$

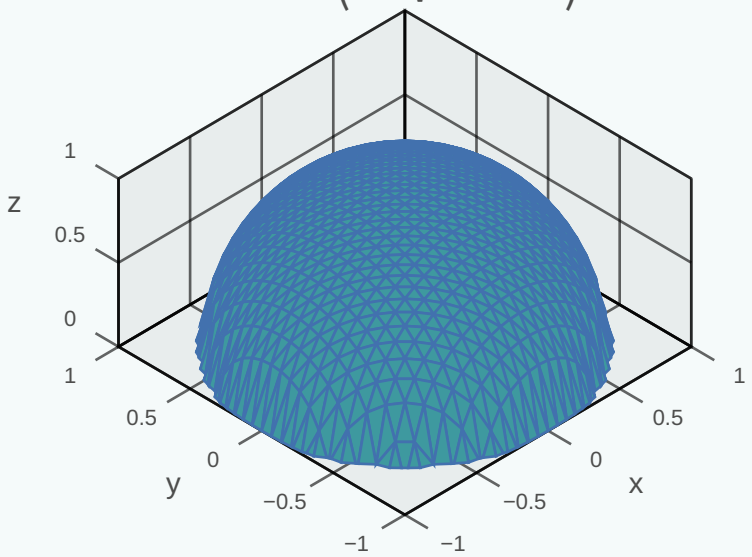

2.3.1. Curves and surfaces

Term	*Definition*
Curves	Vector valued functions of one variable, i.e. they map one-dimensional input to multi-dimensional output
n-dimensional curve	A continuous function $f : I \to Z$ that maps an interval $I \subset \mathbb{R}$ into a subset $Z \subset \mathbb{R}^n$
Surfaces	Real valued functions of multiple variables, i.e. they map multi-dimensional input to one dimensional output
n-dimensional hyper-surface	A continuous real valued function $f : D \to Z$ that maps a sufficiently large subset $D \subset \mathbb{R}^n$ into a subset $Z \subset \mathbb{R}$ If $n = 2$ a hypersurface is also simply called surface. If $n = 1$ a hyper-surface is simply a continuous real valued function of one variable.

The graph of both, an n-dimensional curve and an n-dimensional hyper-surface, is a subset of $\mathbb{R}^{n+1}$. A graphical illustration of such a graph therefore requires $n \leq 2$. However, unlike surfaces which are typically illustrated as in Figure 1 , a curve $c : I \to \mathbb{R}^n$ is usually not illustrated through

$$\operatorname{graph} c = \{(t, y) \in R \times R^n \mid y = c(t) \wedge t \in I\}$$

Instead curves are typically visualized by skipping the independent variable t from the graph, i.e. by visualizing the image of the curve

$$c(I) = \{y \in R^n \mid y = c(t) \wedge t \in I\}$$

which means that we are projecting the graph of the curve onto the gray plane that is spanned from the dependent variables. As this kind of plot saves one dimension we can also visualize curves for which the range Z is a subset of $\mathbb{R}^3$.

TODO:
pt3d bugfixes

***Example** 10*

The graph of the function

$$f : \begin{cases} \mathbb{R} \to \mathbb{R}^2 \\ t \mapsto (\cos t, \sin t) \end{cases}$$

is

$$\begin{aligned} \operatorname{graph}(f) &= \left\{(t, x, y) \in \mathbb{R}^3 \;\middle|\; (x, y) = f(t) \wedge t \in \mathbb{R}\right\} \\ &= \left\{(t, \cos t, \sin t) \in \mathbb{R}^3 \;\middle|\; t \in \mathbb{R}\right\} \end{aligned}$$

and its image is the set

$$\begin{aligned} f(\mathbb{R}) &= \left\{f(t) \in \mathbb{R}^2 \;\middle|\; t \in \mathbb{R}\right\} \\ &= \left\{(\cos t, \sin t) \in \mathbb{R}^2 \mid t \in \mathbb{R}\right\} \end{aligned}$$

TODO:
3d axis order

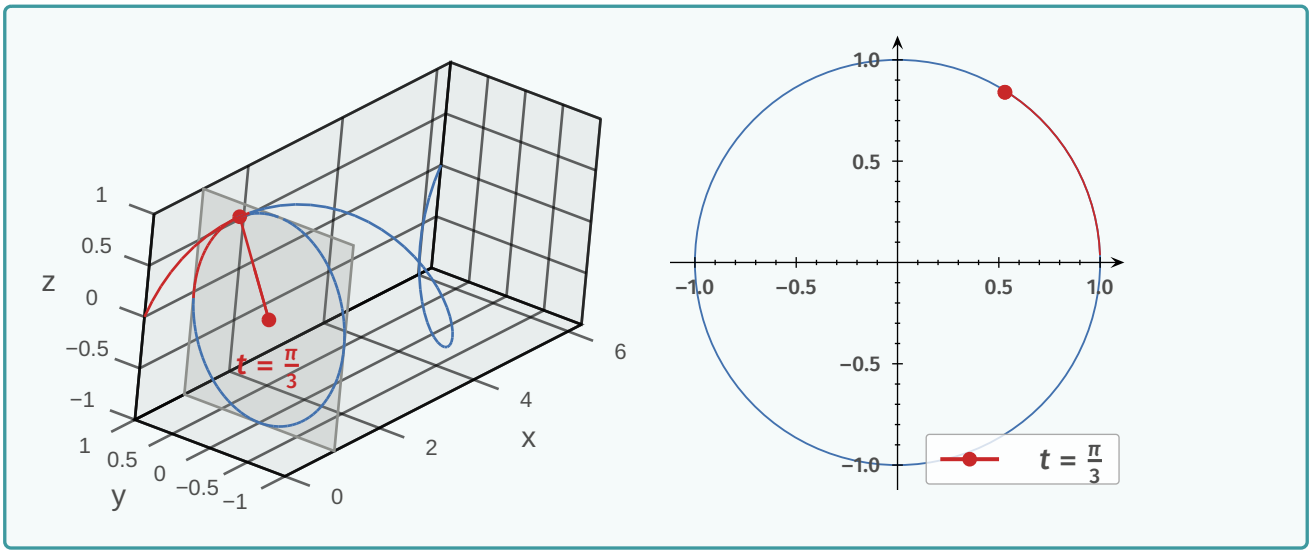

Unlike for image processing, curves are of little importance in machine learning. On the contrary, hypersurfaces belong to the most important functions in machine learning. This is because, to a large extend, machine learning can be thought of being a branch of statistics. In statistics the most important class of functions are probability distributions, which are functions mapping subsets of $\boldsymbol{R}^n$ into $\boldsymbol{R}^+$.

***Term 2* : Histogram**

Consider having some amount of 2D datapoints:

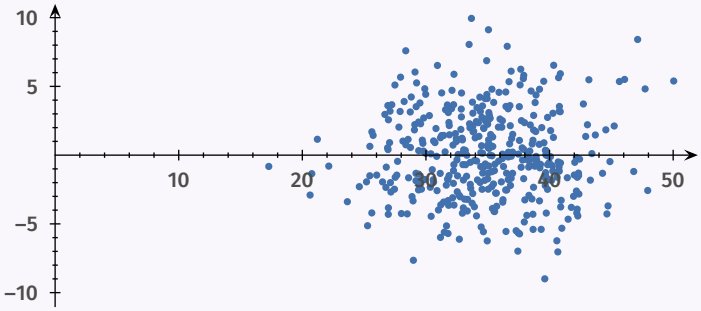

We can place a grid over the diagram and count, how many points are in each of the boxes. We can now plot a diagram that maps the $\boldsymbol{x}$ and $\boldsymbol{y}$ coordinate of the center of each of the boxes to the number of hits of this particular box and thus end up with a diagram that statisticians called **histogram**.

TODO:
proper implementation

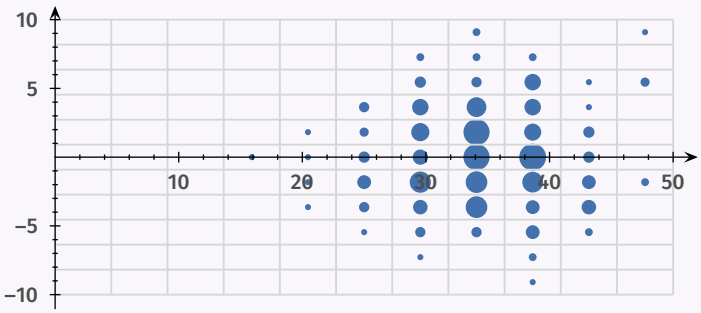

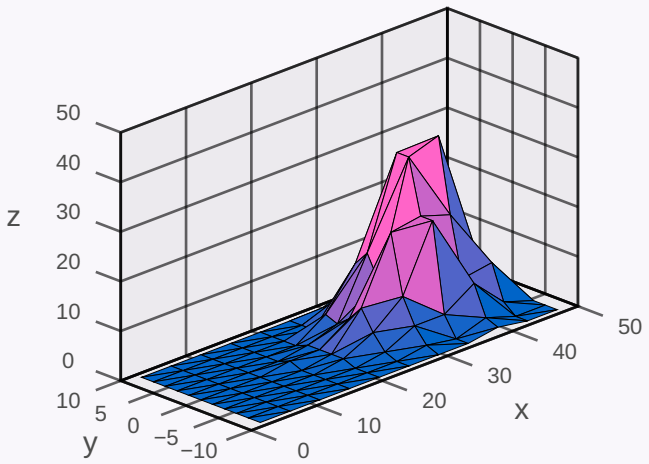

We can try to create a histogram with boxes of infinitely small size. The problem with this approach is, that the number of hits per box decreases when the size of the boxes is getting smaller.

In order to compensate that, we therefore need more datapoints. If we do that, the histogram is being transformed into a function, that associates a number to each position on the diagram, that can be interpreted as likelihood for a point to be almost next to this position: the higher this number is, the more likely it is. This function is called the **probability density** of the experiment.

In addition to probability theory, geometric concepts are used in machine learning. One approach to automatically decide, whether two images are similar or not, is to first extract features (such as corner-points, edges, etc.) from each of the images, and store these features in a list of numbers, which is called **feature vector**. It is reasonable to assume, that similar images are mapped to similar feature vectors.

Now, in vector geometry, there is a concept of similarity: two vectors $\boldsymbol{v}$ and $\boldsymbol{w}$ are similar, if both vectors head into the same direction and are of similar length, so if the difference $\boldsymbol{v} - \boldsymbol{w}$ between both vectors has a small length. In vector geometry, the length of a vector is called a **norm**, and the distance between two vectors is called a **metric**. A norm $\|\cdot\|$ is thus a function that associates with a vector $\boldsymbol{x}$ a positive number $\|\boldsymbol{x}\|$, which is interpreted as length of $\boldsymbol{x}$. One of the most important norms it the Euclidean norm, but there are others. The Euclidean norm is defined as follows:

$$\|\cdot\| : \begin{cases} \mathbb{R}^n \to \mathbb{R}^+ \\ x \quad \mapsto \sqrt{\sum_{i=1}^n x_i^2} \end{cases}$$

= n-dimensional Euclidean norm or l^2-norm, written also as $\|x\|_2$.

Once a vector space is equipped with a norm $\|.\|$, we can also associate a metric with that vector space, which uses the notion of a length to measure the distance between pairs of vectors:

$$d(\cdot,\cdot) : \begin{cases} \mathbb{R}^n \times \mathbb{R}^n \to \mathbb{R}^+ \\ (u,v) \quad \mapsto \|u - v\| \end{cases}$$

= metric. The value $d(u,v)$ is identical to the length of the vector that heads from $\boldsymbol{v}$ to $\boldsymbol{u}$.

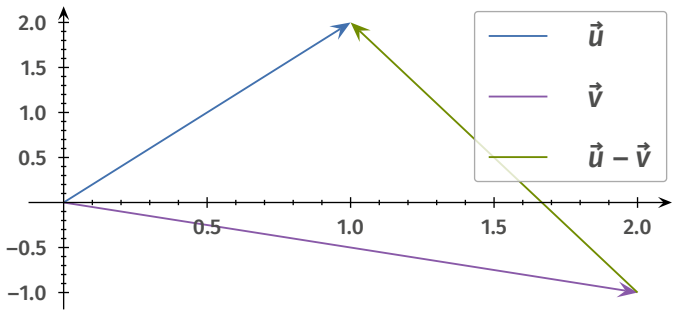

2.3.2. Vector valued functions of several variables

Besides curves and surfaces there are functions that take multi-dimensional input and deliver multi-dimensional output. In order to get an idea how these functions look like, we often display low-dimensional "sections" of the function that can either be displayed as curves or as surfaces. There is however another way, how these functions can be displayed, namely using a **vector field**.

Example 11 **: Spaceship**

A space-ship traveling close to the sun is experiencing a gravitational force that - outside the sun - can be described using the following function:

$$F : \begin{cases} \{r \in \mathbb{R}^3 \mid |r| > r_{\text{Sun}}\} \to \mathbb{R}^3 \\ r \mapsto -G \cdot m_{\text{Sun}} m_{\text{Ship}} \frac{r}{\|r\|^3} \end{cases}$$

r_{Sun} = radius of the sun, G = gravitational constant, $m_{\text{Sun}}, m_{\text{Ship}}$ = masses of the sun and space ship.

If the center of the sun is located at $r = 0$, the input vector r denotes the position of the space ship as seen from the center of the sun and F denotes the gravitational force that the space ship is experiencing.

From a mathematical perspective the constants are not important. Therefore we simplify the formula by setting all of the constants to 1, which leads to:

$$F : \begin{cases} \{r \in \mathbb{R}^3 \mid |r| > 1\} \to \mathbb{R}^3 \\ r \mapsto -\frac{r}{\|r\|^3} \end{cases}$$

In order to better understand the formula of the force field, we calculate the force experienced from the space ship at position $r = (1, 2, 2)^T$. Note first, that r is in the domain of definition of F , because its norm is greater than $r_{\text{Sun}} = 1$

$$\|r\| = \sqrt{1^2 + 2^2 + 2^2} = \sqrt{9} = 3$$

This simply indicates, that the spaceship resides at a location exterior to the sun. We can therefore use this to calculate the force that the space ship is experiencing:

$$F\begin{pmatrix}1\\2\\2\end{pmatrix} = -\frac{1}{\|(1,2,2)^T\|^3}\begin{pmatrix}1\\2\\2\end{pmatrix} = -\frac{1}{3^3}\begin{pmatrix}1\\2\\2\end{pmatrix} = -\frac{1}{27}\begin{pmatrix}1\\2\\2\end{pmatrix}$$

We can now continue to calculate values of the force F at other locations and try to graphically visualize the result. Since the graph of F

$$\text{graph}(F) = \left\{ (r, f) \in \mathbb{R}^3 \times \mathbb{R}^3 \;\middle|\; \|r\| > 1 \wedge f = -\frac{r}{\|r\|^3} \right\}$$

is a subset of a $3 + 3 = 6$-dimensional space, we cannot however not draw the graph of F directly in any meaningful way. A possibility to still visualize this function, is to draw the function as **vector field**. By this we mean that we place 3-dimensional arrows representing F at different points in a 3-dimensional space that identify the physical location of the space ship. The arrow that is placed at position r has a length proportional to the norm of the force $F(r)$ and a direction parallel to the direction of the force. It thus tells you the strength and direction in which the space ship is pulled due to the gravitational force.

TODO:
pt3d vector field

TODO:
2d repr

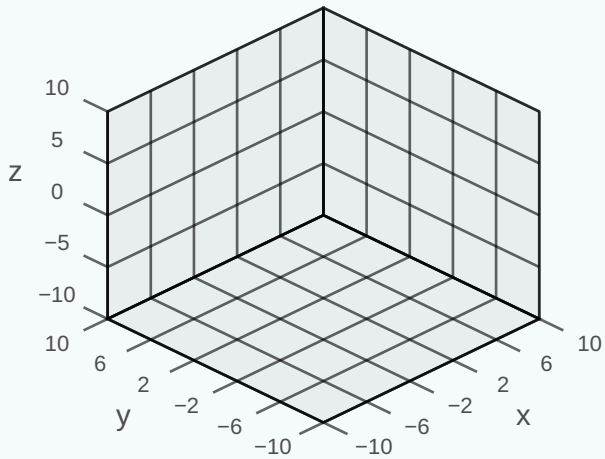

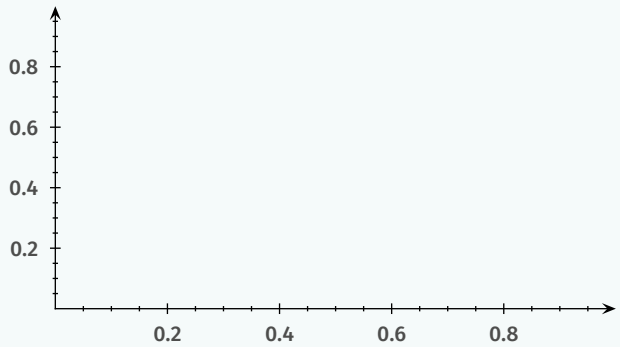

The mathematical tool for determining the “strength of the field” is the **norm of a vector**, which measures the length of a vector and thus encodes the strength of a field. If we calculate $\|F(r)\|$ in our case, we get

$$\|F(r)\| = \left\| -\frac{r}{\|r\|^3} \right\| = \frac{\|r\|}{\|r\|^3} = \frac{1}{\|r\|^2}$$

which illustrates that the force is getting smaller, the larger $|r|$ is

The mathematical tool for determining the “direction of a field” are **unit vectors**. The unit vector of r is a vector of length one, that points into the same direction as r. It thus encodes information regarding the direction of r, but not regarding its length. The unit vector of r is usually denoted by e_r. It can be calculated via

$$e_r = \frac{1}{\|r\|} \cdot r$$

In machine learning many functions are vector valued functions of many variables. A very important example that falls into this category are linear transformations. Linear transformations are of utmost importance in image processing, as you need them to calculate the camera-perspective of any given scene. Furthermore, neural networks are typically defined to be a stack of unknown linear transformations, each of which being followed by a rather simple, predefined non-linear operation. Training a neural network is therefore almost equivalent to finding suitable linear transformations that make the network perform the desired operation.

2.4. CALCULUS OF CURVES

***Term 3* : Coordinate functions**

Let $I \subset \mathbb{R}$ be an interval and $c : I \to \mathbb{R}^n$ be an arbitrary curve. Then there are n (contiguous) coordinate functions $c_i : I \to \mathbb{R}$, such that

$$c(t) = \begin{pmatrix} c_1(t) \\ c_2(t) \\ \vdots \\ c_n(t) \end{pmatrix}$$

and vice versa: If n real valued (continuous) functions $c_i : I \to \mathbb{R}$ with a common interval I as domain of definition are given, we can define a curve $c(t)$ whose coordinate functions are $c_1, c_2, \ldots, c_n$.

The decomposition of curves into coordinate functions allows us to generalize most of the concepts that you have learned in prior analysis courses to curves: We can for instance define limit points, derivatives and Taylor series of curves by applying the corresponding concept of real valued functions to all of the coordinate functions. The remaining challenge then is to interpret the results.

***Example 12* : nyah :3 i have heart issues caused by a constant state of distress and extreme substance abuse in my youth UwU**

We consider the curve

$$f : \begin{cases} (0;\infty) \to \mathbb{R}^2 \\ t \qquad \mapsto \begin{pmatrix} \cos(t) \\ \frac{\sin(t)}{t} \end{pmatrix} \end{cases}$$

The component functions are

$$f_1(t) : \cos(t)$$
$$f_2(t) : \frac{\sin(t)}{t}$$

We can thus argue:

Continuity

Since both of the component functions are continuous, f is continuous.

Limit point

$$\lim_{t\to 0} f_1(t) = \lim_{t\to 0} \cos(t) = \cos(0) = 1$$

$$\lim_{t\to 0} f_2(t) = \lim_{t\to 0} \frac{\sin(t)}{t} \overset{\text{"0/0"}}{=} \lim_{t\to 0} \frac{\frac{d}{dt}\sin(t)}{\frac{d}{dt}t} = \lim_{t\to 0} \frac{\cos(t)}{1} = \cos(0) = 1$$

$$\Rightarrow \lim_{t\to 0} f(t) = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

Derivative

$$\frac{d}{dt}f_1(t) = \frac{d}{dt}\cos(t) = -\sin(t)$$
$$\frac{d}{dt}f_2(t) = \frac{d}{dt}\frac{\sin(t)}{t} = \frac{\cos(t)\cdot t - \sin(t)\cdot 1}{t^2} = \frac{\cos(t)}{t} - \frac{\sin(t)}{t^2}$$
$$\Rightarrow \frac{d}{dt}f(t) = \begin{pmatrix} -\sin(t) \\ \frac{\cos(t)}{t} - \frac{\sin(t)}{t^2} \end{pmatrix}$$

Linearisation at $t = \pi$

$$f_1(t) \approx f_1(\pi) + f_1'(\pi)(t-\pi) = \cos(\pi) - \sin(\pi)(t-\pi) = -1$$
$$f_2(t) \approx f_2(\pi) + f_2'(\pi)(t-\pi) = \frac{\sin(\pi)}{\pi} + \left(\frac{\cos(\pi)}{\pi} - \frac{\sin(\pi)}{\pi^2}\right)(t-\pi) = -\frac{1}{\pi}(t-\pi)$$
$$\Rightarrow f(t) \approx f(\pi) + f'(\pi)(x-\pi) = \begin{pmatrix} -1 \\ -\frac{1}{\pi}(t-\pi) \end{pmatrix} = \begin{pmatrix} -1 \\ 0 \end{pmatrix} - (t-\pi)\begin{pmatrix} 0 \\ \frac{1}{\pi} \end{pmatrix}$$

TODO:
diagram

For any $\vec{a}, \vec{v} \in \mathbb{R}^n$ a curve of the form $g : \begin{cases} \mathbb{R} \to \mathbb{R}^n \\ t \mapsto \vec{a} + t \cdot \vec{v} \end{cases}$ corresponds to a line going through the point a in the direction of v.

TODO:
Notes 30/31

Term 4 : **Reparametrization**

Let

$$c : \begin{cases} I \to \mathbb{R}^n \\ t \mapsto c(t) \end{cases}$$

be a curve and $h : J \to I$ an arbitrary bijective function. Then the image of the curve

$$d : \begin{cases} J \to \mathbb{R}^n \\ s \mapsto c(h(s)) \end{cases}$$

is identical to the image of the curve $c(t)$:

$$c(I) = d(J)$$

We call $d(s)$ ***reparametrization*** of $c(t)$ using $t = h(s)$

TODO:
rest of chapter

Term 5 : **Image of a curve**

The image of a curve

$$c : \begin{cases} I \to \mathbb{R}^n \\ t \mapsto c(t) \end{cases}$$

can be envisioned as a road that is passed from a passenger, who reaches position $c(t)$ at time t. The derivative $c'(t)$ is a measure of the speed of the person at time t. This implies that the derivative $c'(t)$ points into a direction tangent to the road at position $c(t)$ and that the norm $\|c'(t)\|$ of the derivative is a measure of the speed with which the person is moving at time t.

If we reparametrize the curve c with a bijective function $t = h(s)$ the image of the curve

$$d(s) = c(h(s))$$

is identical to the image of c and thus corresponds to the same road. It is however passed by another person, who is passing point $d(s) = c(h(s))$ with a velocity

$$d'(s) = c'(h(s)) \cdot h'(s)$$

that is $h'(s)$ faster (or slower) than the velocity of the first person at time $t = h(s)$.

TODO:
formal proof (Notes 33/34)

2.5. CALCULUS OF REAL VALUED FUNCTIONS IN MANY VARIABLES